

Data analysis and visualization

```
In [1]: %matplotlib inline
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import os
from mpl_toolkits.basemap import Basemap
```

```
In [2]: # Input data from "/data" director
os.listdir("data")
```

```
Out[2]: ['label_categories.csv',
'app_labels.csv',
'sample_submission.csv',
'phone_brand_device_model.csv',
'events.csv',
'gender_age_train.csv',
'app_events.csv',
'gender_age_test.csv']
```

```
In [3]: import seaborn as sns
sns.set(color_codes=True)
app_event=pd.read_csv("data/events.csv")
app_event.shape
```


```
Out[3]: (3252950, 5)
```

```
In [4]: app_event.timestamp=pd.to_datetime(app_event.timestamp)
app_event['time_hour'] = app_event.timestamp.apply(lambda x: x.hour)
```


```
In [5]: # Show frequency of events by hour  
app_event['time_hour'].value_counts()
```

```
Out[5]: 10    183839  
        21    181175  
        20    178179  
        22    168246  
        11    167025  
        19    166160  
         9    166061  
        12    162745  
         8    157896  
        18    156209  
        13    155337  
        17    153516  
        14    151379  
        16    150732  
        15    149912  
         7    146667  
        23    136339  
         0    120512  
         6    116370  
         1     72671  
         5     66411  
         2     53764  
         3     47048  
         4     44757  
Name: time_hour, dtype: int64
```

```
In [6]: ax = sns.countplot(x="time_hour", data=app_event)
```




```
In [7]: import calendar
app_event['week_day'] = app_event.timestamp.apply(lambda x: calendar.day_name[x.weekday()])
ax = sns.countplot(x="week_day", data=app_event)
```




```
In [8]: gender=pd.read_csv("data/gender_age_train.csv")
print(gender.gender.value_counts())
ax = sns.countplot(x="gender", data=gender)
```


```
M    47904
F    26741
Name: gender, dtype: int64
```



```
In [9]: # Distribution by age  
sns.distplot(gender.age, hist=False);
```




```
In [10]: # Distribution by sex  
sns.kdeplot(gender.age[gender.gender=="M"], label="Male")  
sns.kdeplot(gender.age[gender.gender=="F"], label="Female")  
plt.legend();
```



Female at old age are using mobiles little bit more than males at old age


```
In [11]: print("Male mobile usage count by age")
ax = sns.countplot(x="group", data=gender[gender.gender=="M"])
```

Male mobile usage count by age



```
In [12]: print("Female mobile usage count by age")
ax = sns.countplot(x="group", data=gender[gender.gender=="F"])
```

Female mobile usage count by age




```
In [13]: appscategories=pd.read_csv("data/label_categories.csv")
print(appscategories.head())
print(appscategories.shape)
```

```
   label_id  category
0         1         NaN
1         2  game-game type
2         3  game-Game themes
3         4  game-Art Style
4         5  game-Leisure time
(930, 2)
```

Joint visualisation - male and female

```
In [14]: gender.group.value_counts().sort_index(ascending=False).plot(kind='barh')
```


```
Out[14]: <matplotlib.axes._subplots.AxesSubplot at 0x7f8b64c0ceb8>
```



Further break down by age

```
In [15]: c = gender.groupby(['age', 'gender']).size().unstack().reindex(index=np.arange(gender.age.min(), gender.age.max()+1)).fillna(0)
ax1, ax2 = c.plot(kind='bar', figsize=(22,6), subplots=True);
ax1.vlines(np.array([23,26,28,32,42])-0.5,0,1800,alpha=0.5,linewidth=1,color='r')
ax2.vlines(np.array([22,26,28,31,38])-0.5,0,3000,alpha=0.5,linewidth=1,color='r')
```

```
Out[15]: <matplotlib.collections.LineCollection at 0x7f8b65467668>
```



Locations visualization

```
In [16]: df_events = pd.read_csv("data/events.csv", dtype={'device_id': np.str})
df_events.head()
```

Out[16]:

	event_id	device_id	timestamp	longitude	latitude
0	1	29182687948017175	2016-05-01 00:55:25	121.38	31.24
1	2	-6401643145415154744	2016-05-01 00:54:12	103.65	30.97
2	3	-4833982096941402721	2016-05-01 00:08:05	106.60	29.70
3	4	-6815121365017318426	2016-05-01 00:06:40	104.27	23.28
4	5	-5373797595892518570	2016-05-01 00:07:18	115.88	28.66

Locations of events - World map

Expand size of the screen so larger map can nicely fit without truncated window with scroller.

```
In [17]: %%html
<style>
.output_wrapper, .output {
    height:auto !important;
    max-height:1000px; /* your desired max-height here */
}
.output_scroll {
    box-shadow:none !important;
    webkit-box-shadow:none !important;
}
</style>
```

```
In [18]: # Set plot
df_events_sample = df_events.sample(n=90000)
plt.figure(1, figsize=(20,10))
pd.set_option('display.max_colwidth', -1)

# Map of World
map = Basemap(projection='merc',
              llcrnrlat=-60,
              urcrnrlat=65,
              llcrnrlon=-120,
              urcrnrlon=180,
              lat_ts=0,
              resolution='c')

map.fillcontinents(color='#500000',lake_color='#000000') # grey land, black
Lakes
map.drawmapboundary(fill_color='#202020')                # black background
map.drawcountries(linewidth=0.1, color="w")              # white line of co
untry borders

# Plot the data
mxy = map(df_events_sample["longitude"].tolist(), df_events_sample["latitud
e"].tolist())
map.scatter(mxy[0], mxy[1], s=3, c="#12AABB", lw=0, alpha=1, zorder=5)

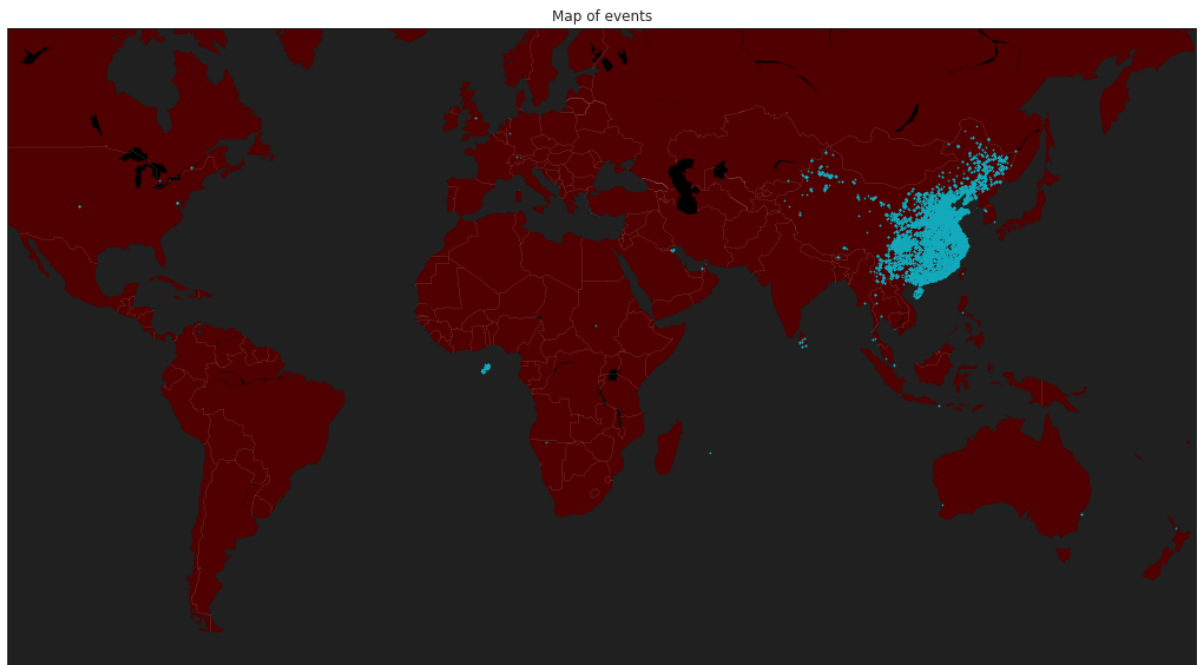
plt.title("Map of events")
plt.show()
```



```

/root/anaconda3/lib/python3.6/site-packages/mpl_toolkits/basemap/__init__.p
y:1767: MatplotlibDeprecationWarning: The get_axis_bgcolor function was dep
recated in version 2.0. Use get_facecolor instead.
    axisbgc = ax.get_axis_bgcolor()
/root/anaconda3/lib/python3.6/site-packages/mpl_toolkits/basemap/__init__.p
y:1698: MatplotlibDeprecationWarning: The axesPatch function was deprecated
in version 2.1. Use Axes.patch instead.
    limb = ax.axesPatch
/root/anaconda3/lib/python3.6/site-packages/mpl_toolkits/basemap/__init__.p
y:3222: MatplotlibDeprecationWarning: The ishold function was deprecated in
version 2.0.
    b = ax.ishold()
/root/anaconda3/lib/python3.6/site-packages/mpl_toolkits/basemap/__init__.p
y:3231: MatplotlibDeprecationWarning: axes.hold is deprecated.
    See the API Changes document (http://matplotlib.org/api/api\_changes.htm
l)
    for more details.
    ax.hold(b)

```



As we can see some events have $(lat,lon) = (0,0)$ which probably means that location couldn't be determined. We can find all event on that location and all events that have longitud and latitud less than 1 which means they are really close to that location. That is location is not probable since it's points to sea area close to African cost.

```
In [19]: df_at0 = df_events[(df_events["longitude"]==0) & (df_events["latitude"]==0
)]
df_near0 = df_events[(df_events["longitude"]>-1) &\
                      (df_events["longitude"]<1) &\
                      (df_events["latitude"]>-1) &\
                      (df_events["latitude"]<1)]

print("Total number of events:", len(df_events))
print("Number of events at (0,0):", len(df_at0))
print("Number of events near (0,0):", len(df_near0))
```

```
Total number of events: 3252950
Number of events at (0,0): 968675
Number of events near (0,0): 969871
```

Locations of events - zooming in region of China

```
In [20]: # Locate China region
lon_min, lon_max = 75, 135
lat_min, lat_max = 15, 55

idx_china = (df_events["longitude"]>lon_min) &\
            (df_events["longitude"]<lon_max) &\
            (df_events["latitude"]>lat_min) &\
            (df_events["latitude"]<lat_max)

df_events_china = df_events[idx_china].sample(n=100000)

# China
plt.figure(2, figsize=(20,15))

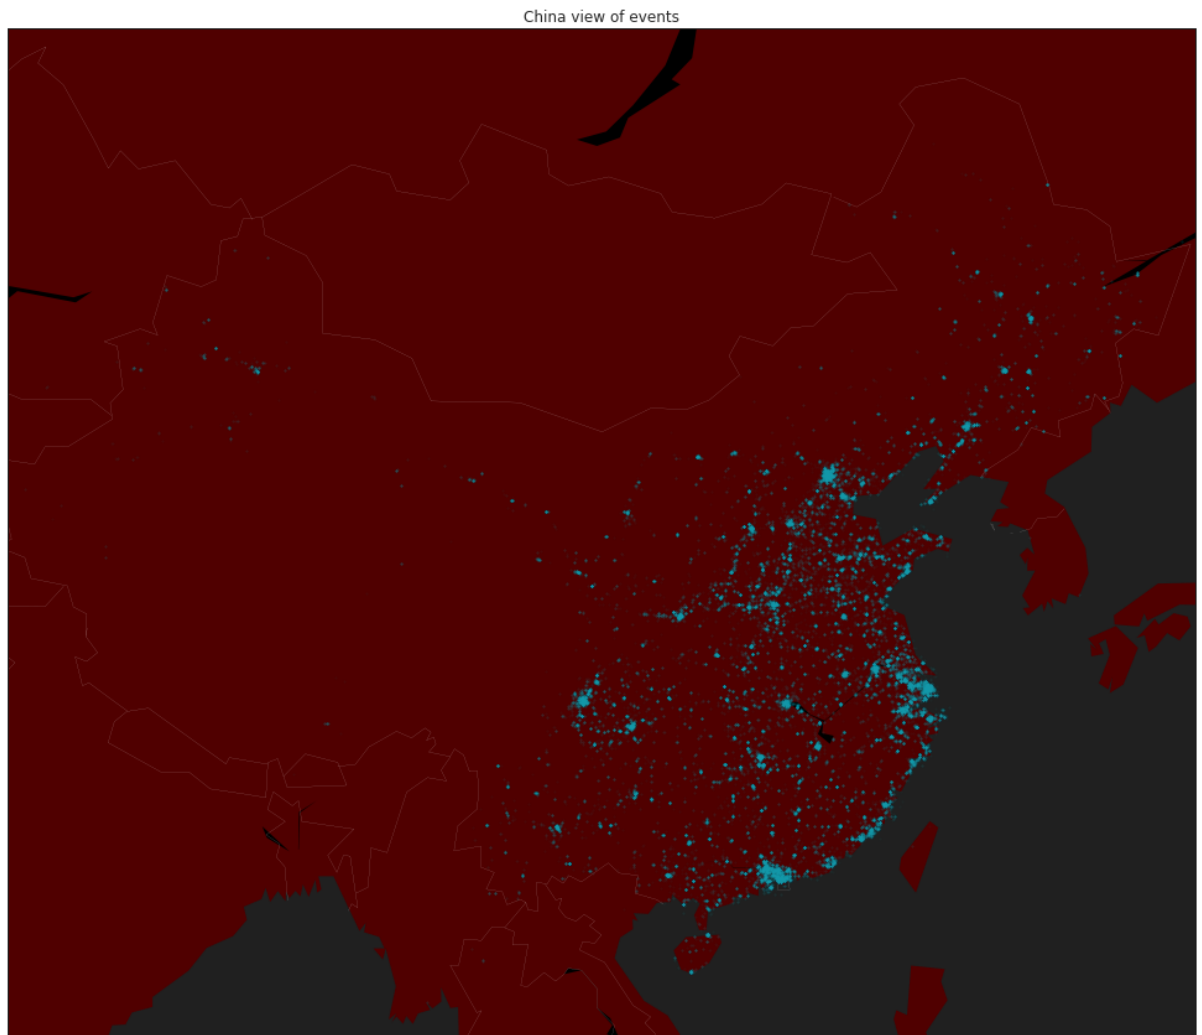
map_zoom = Basemap(projection='merc',
                    llcrnrlat=lat_min,
                    urcrnrlat=lat_max,
                    llcrnrlon=lon_min,
                    urcrnrlon=lon_max,
                    lat_ts=35,
                    resolution='c')

map_zoom.fillcontinents(color='#500000',lake_color='#000000') # dark grey land, black lakes
map_zoom.drawmapboundary(fill_color='#202020')                # black background
map_zoom.drawcountries(linewidth=0.1, color="w")              # thin white line for country borders

# Plot the data
mxy = map_zoom(df_events_china["longitude"].tolist(), df_events_china["latitude"].tolist())
map_zoom.scatter(mxy[0], mxy[1], s=5, c="#12AABB", lw=0, alpha=0.05, zorder=5)

plt.title("China view of events")
plt.show()
```

```
/root/anaconda3/lib/python3.6/site-packages/mpl_toolkits/basemap/__init__.p
y:1767: MatplotlibDeprecationWarning: The get_axis_bgcolor function was dep
recated in version 2.0. Use get_facecolor instead.
    axisbgc = ax.get_axis_bgcolor()
/root/anaconda3/lib/python3.6/site-packages/mpl_toolkits/basemap/__init__.p
y:1698: MatplotlibDeprecationWarning: The axesPatch function was deprecated
in version 2.1. Use Axes.patch instead.
    limb = ax.axesPatch
/root/anaconda3/lib/python3.6/site-packages/mpl_toolkits/basemap/__init__.p
y:3222: MatplotlibDeprecationWarning: The ishold function was deprecated in
version 2.0.
    b = ax.ishold()
/root/anaconda3/lib/python3.6/site-packages/mpl_toolkits/basemap/__init__.p
y:3231: MatplotlibDeprecationWarning: axes.hold is deprecated.
    See the API Changes document (http://matplotlib.org/api/api\_changes.htm
l)
    for more details.
    ax.hold(b)
```



This map nicely shows population density of China. We'll analyze one city region in this case we'll take Shanghai. What follows are maps that are showing longitude and latitude areas. In same way, we can zoom on any area on earth for showing more details in that area.

Locations of applications events - region of Shanghai

```
In [21]: # Shanghai sits on 31.2304° N, 121.4737° E
# Sampling wider Shanghai region
lon_min, lon_max = 115, 125
lat_min, lat_max = 28, 35

idx_shanghai = (df_events["longitude"]>lon_min) &\
                (df_events["longitude"]<lon_max) &\
                (df_events["latitude"]>lat_min) &\
                (df_events["latitude"]<lat_max)

df_events_shanghai = df_events[idx_shanghai]

# Map of Shanghai region
plt.figure(3, figsize=(20,15))

m_shanghai = Basemap(projection='merc',
                      llcrnrlat=lat_min,
                      urcrnrlat=lat_max,
                      llcrnrlon=lon_min,
                      urcrnrlon=lon_max,
                      lat_ts=35,
                      resolution='c')

m_shanghai.fillcontinents(color='#500000',lake_color='#000000') # dark land, black lakes
m_shanghai.drawmapboundary(fill_color='#000000') # black background
m_shanghai.drawcountries(linewidth=0.1, color="w") # white line for country borders

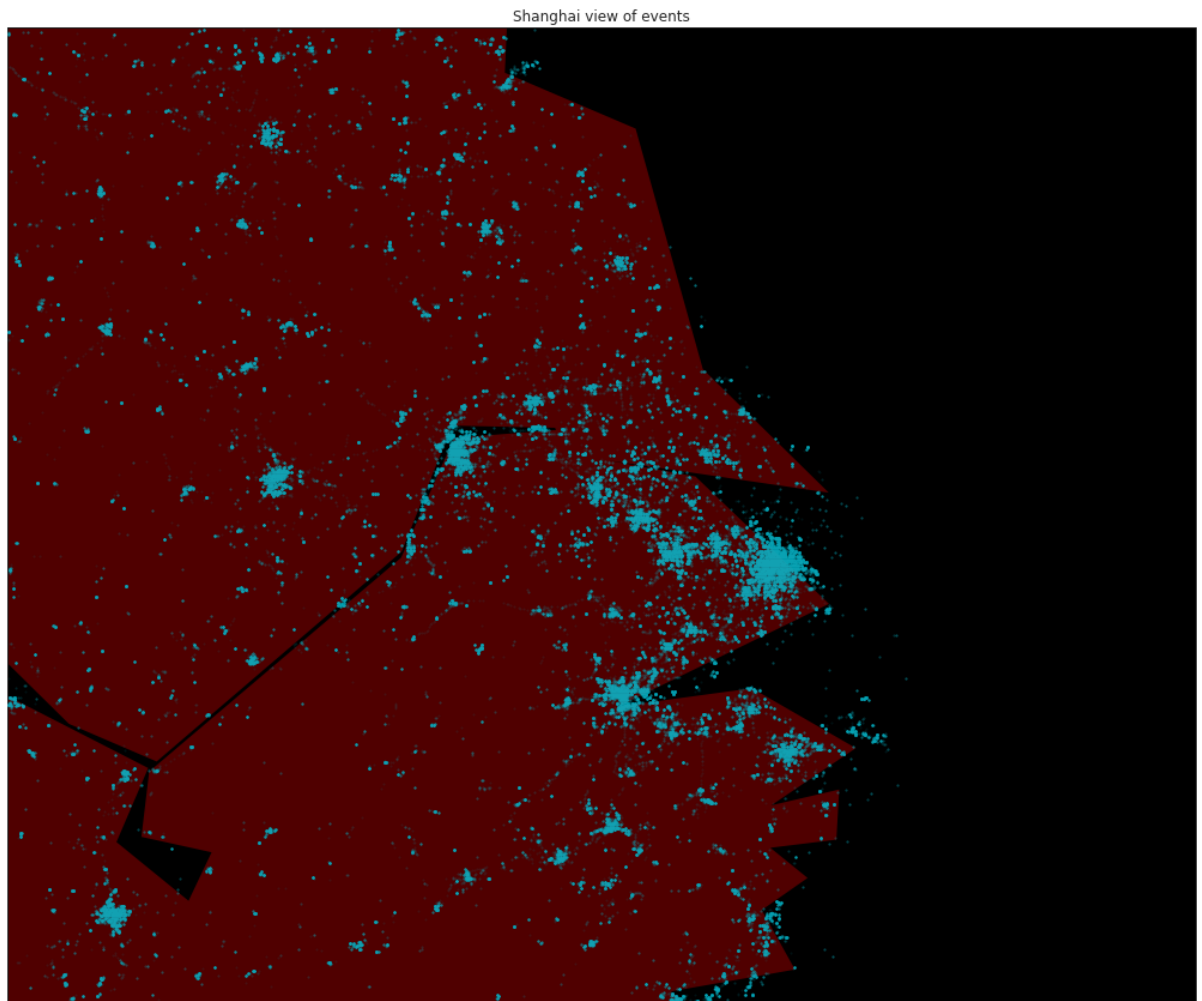
# Plot the data
mxy = m_shanghai(df_events_shanghai["longitude"].tolist(), df_events_shanghai["latitude"].tolist())
m_shanghai.scatter(mxy[0], mxy[1], s=5, c="#12AABB", lw=0, alpha=0.1, zorder=5)

plt.title("Shanghai view of events")
plt.show()
```

```

/root/anaconda3/lib/python3.6/site-packages/mpl_toolkits/basemap/__init__.p
y:1767: MatplotlibDeprecationWarning: The get_axis_bgcolor function was dep
recated in version 2.0. Use get_facecolor instead.
    axisbgc = ax.get_axis_bgcolor()
/root/anaconda3/lib/python3.6/site-packages/mpl_toolkits/basemap/__init__.p
y:1698: MatplotlibDeprecationWarning: The axesPatch function was deprecated
in version 2.1. Use Axes.patch instead.
    limb = ax.axesPatch
/root/anaconda3/lib/python3.6/site-packages/mpl_toolkits/basemap/__init__.p
y:3222: MatplotlibDeprecationWarning: The ishold function was deprecated in
version 2.0.
    b = ax.ishold()
/root/anaconda3/lib/python3.6/site-packages/mpl_toolkits/basemap/__init__.p
y:3231: MatplotlibDeprecationWarning: axes.hold is deprecated.
    See the API Changes document (http://matplotlib.org/api/api\_changes.htm
1)
    for more details.
    ax.hold(b)

```



We can see that population around big cities is very dispersed.

Now we'll show male and female app events

Male and female app events in region of Shanghai

```
In [22]: # Load the train data and join on the events
df_train = pd.read_csv("data/gender_age_train.csv", dtype={'device_id': np.str})

df_plot = pd.merge(df_train, df_events_shanghai, on="device_id", how="inner")

df_m = df_plot[df_plot["gender"]=="M"]
df_f = df_plot[df_plot["gender"]=="F"]
```

Visualize male and female events in Shanghai


```

In [23]: # Female and male plot

plt.figure(4, figsize=(20,10))

plt.subplot(121)
m_sh_m = Basemap(projection='merc',
                  llcrnrlat=lat_min,
                  urcrnrlat=lat_max,
                  llcrnrlon=lon_min,
                  urcrnrlon=lon_max,
                  lat_ts=35,
                  resolution='c')
m_sh_m.fillcontinents(color='#500000',lake_color='#000000') # dark grey land, black lakes
m_sh_m.drawmapboundary(fill_color='#000000') # black background
m_sh_m.drawcountries(linewidth=0.1, color="w") # thin white line for country borders
mxy = m_sh_m(df_m["longitude"].tolist(), df_m["latitude"].tolist())
m_sh_m.scatter(mxy[0], mxy[1], s=5, c="#1292db", lw=0, alpha=0.1, zorder=5)
plt.title("Male events in Shanghai")

plt.subplot(122)
m_sh_f = Basemap(projection='merc',
                  llcrnrlat=lat_min,
                  urcrnrlat=lat_max,
                  llcrnrlon=lon_min,
                  urcrnrlon=lon_max,
                  lat_ts=35,
                  resolution='c')
m_sh_f.fillcontinents(color='#500000',lake_color='#000000') # dark grey land, black lakes
m_sh_f.drawmapboundary(fill_color='#000000') # black background
m_sh_f.drawcountries(linewidth=0.1, color="w") # thin white line for country borders
mxy = m_sh_f(df_f["longitude"].tolist(), df_f["latitude"].tolist())
m_sh_f.scatter(mxy[0], mxy[1], s=5, c="#fd3096", lw=0, alpha=0.1, zorder=5)
plt.title("Female events in Shanghai")

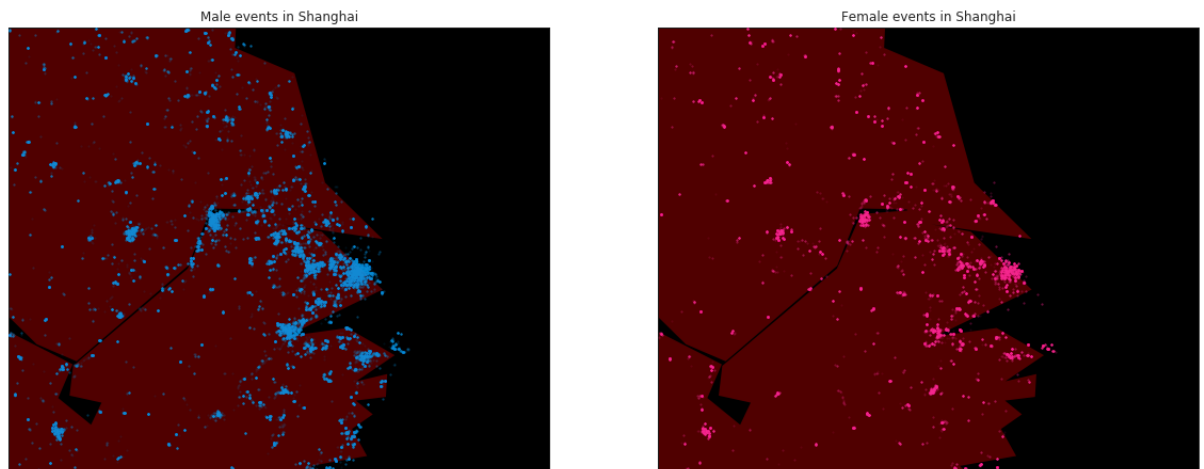
plt.show()

```

```

/root/anaconda3/lib/python3.6/site-packages/mpl_toolkits/basemap/__init__.p
y:1767: MatplotlibDeprecationWarning: The get_axis_bgcolor function was dep
recated in version 2.0. Use get_facecolor instead.
    axisbgc = ax.get_axis_bgcolor()
/root/anaconda3/lib/python3.6/site-packages/mpl_toolkits/basemap/__init__.p
y:1698: MatplotlibDeprecationWarning: The axesPatch function was deprecated
in version 2.1. Use Axes.patch instead.
    limb = ax.axesPatch
/root/anaconda3/lib/python3.6/site-packages/mpl_toolkits/basemap/__init__.p
y:3222: MatplotlibDeprecationWarning: The ishold function was deprecated in
version 2.0.
    b = ax.ishold()
/root/anaconda3/lib/python3.6/site-packages/mpl_toolkits/basemap/__init__.p
y:3231: MatplotlibDeprecationWarning: axes.hold is deprecated.
    See the API Changes document (http://matplotlib.org/api/api\_changes.htm
l)
    for more details.
    ax.hold(b)

```



For marketing analysis, this might be interesting for further exploration. Which city areas are showing more men activities and which are showing more female activities and in which times of day?

Analysis

Problem classification

Our task is to build a model predicting users' demographic characteristics based on their app usage, geolocation, and mobile device properties. So we need to solve multiclass classification problem. This is a case where one label needs to be predicted based on several others.

Logistic regression

Logistic regression algorithm could be obvious choice for that. In the multiclass case, the training algorithm uses the one-vs-rest (OvR) scheme.

```
In [24]: from sklearn.preprocessing import LabelEncoder
from scipy.sparse import csr_matrix, hstack
from sklearn.linear_model import LogisticRegression
from sklearn.cross_validation import StratifiedKFold
from sklearn.metrics import log_loss
```

```
/root/anaconda3/lib/python3.6/site-packages/sklearn/cross_validation.py:41:
DeprecationWarning: This module was deprecated in version 0.18 in favor of
the model_selection module into which all the refactored classes and functi
ons are moved. Also note that the interface of the new CV iterators are dif
ferent from that of this module. This module will be removed in 0.20.
    "This module will be removed in 0.20.", DeprecationWarning)
```

Loading of data

```
In [25]: # Read gender Train and Test data
datadir = 'data'
g_a_train = pd.read_csv(os.path.join(datadir, 'gender_age_train.csv'), index_
col='device_id')
g_a_test = pd.read_csv(os.path.join(datadir, 'gender_age_test.csv'), index_co
l = 'device_id')
phone = pd.read_csv(os.path.join(datadir, 'phone_brand_device_model.csv'))

print(g_a_train.head())
print("-----")
print(g_a_test.head())
print("-----")
print(phone.head())
```

```

              gender  age  group
device_id
-8076087639492063270  M    35  M32-38
-2897161552818060146  M    35  M32-38
-8260683887967679142  M    35  M32-38
-4938849341048082022  M    30  M29-31
 245133531816851882  M    30  M29-31
-----
Empty DataFrame
Columns: []
Index: [1002079943728939269, -1547860181818787117, 7374582448058474277, -62
20210354783429585, -5893464122623104785]
-----
              device_id  phone_brand  device_model
0 -8890648629457979026   小米         红米
1  1277779817574759137   小米         MI 2
2  5137427614288105724   三星        Galaxy S4
3  3669464369358936369  SUGAR      时尚手机
4 -5019277647504317457   三星        Galaxy Note 2
```

```
In [26]: # Remove duplicate device ids in the phones
phone = phone.drop_duplicates(subset=['device_id'], keep='first').set_index('device_id')
events = pd.read_csv('data/events.csv', parse_dates=['timestamp'], index_col='event_id')
appevents = pd.read_csv('data/app_events.csv', usecols=['event_id', 'app_id', 'is_active'], dtype={'is_active': bool})
applabels = pd.read_csv('data/app_labels.csv')
```

```
/root/anaconda3/lib/python3.6/site-packages/numpy/lib/arraysetops.py:463: FutureWarning: elementwise comparison failed; returning scalar instead, but in the future will perform elementwise comparison
  mask |= (ar1 == a)
```

```
In [27]: print(phone.head())
print("-----")
print(events.head())
print("-----")
print(appevents.head())
print("-----")
print(applabels.head())
print("-----")
```

	phone_brand	device_model
device_id		
-8890648629457979026	小米	红米
1277779817574759137	小米	MI 2
5137427614288105724	三星	Galaxy S4
3669464369358936369	SUGAR	时尚手机
-5019277647504317457	三星	Galaxy Note 2

	device_id	timestamp	longitude	latitude
event_id				
1	29182687948017175	2016-05-01 00:55:25	121.38	31.24
2	-6401643145415154744	2016-05-01 00:54:12	103.65	30.97
3	-4833982096941402721	2016-05-01 00:08:05	106.60	29.70
4	-6815121365017318426	2016-05-01 00:06:40	104.27	23.28
5	-5373797595892518570	2016-05-01 00:07:18	115.88	28.66

	event_id	app_id	is_active
0	2	5927333115845830913	True
1	2	-5720078949152207372	False
2	2	-1633887856876571208	False
3	2	-653184325010919369	True
4	2	8693964245073640147	True

	app_id	label_id
0	7324884708820027918	251
1	-4494216993218550286	251
2	6058196446775239644	406
3	6058196446775239644	407
4	8694625920731541625	406

Main feature selection

Main features chosen are:

- phone brand
- device model
- installed apps
- app labels

We need to one-hot encode everything and put in sparse matrices which will help deal with a very large number of features. Regarding "Phone brand" feature; we'll make two columns that show which train or test set row a particular device_id belongs to.

```
In [28]: g_a_train['trainr'] = np.arange(g_a_train.shape[0])
g_a_test['testr'] = np.arange(g_a_test.shape[0])
```

```
In [29]: print(g_a_train.head())
print("-----")
print(g_a_test.head())
```

device_id	gender	age	group	trainr
-8076087639492063270	M	35	M32-38	0
-2897161552818060146	M	35	M32-38	1
-8260683887967679142	M	35	M32-38	2
-4938849341048082022	M	30	M29-31	3
245133531816851882	M	30	M29-31	4

device_id	testr
1002079943728939269	0
-1547860181818787117	1
7374582448058474277	2
-6220210354783429585	3
-5893464122623104785	4

Constructing sparse matrix of features in following way:

`csr_matrix((data, (row_ind, col_ind)), [shape=(M, N)])` where data, row_ind and col_ind satisfy the relationship `a[row_ind[k], col_ind[k]] = data[k]`

This allows us to define what values to put into certain places in a sparse matrix. For phone brand data the data array will be all ones, row_ind will be the row number of a device and col_ind will be the number of brand.

Brand features

```
In [30]: brand_encoder = LabelEncoder().fit(phone.phone_brand)
phone['brand'] = brand_encoder.transform(phone['phone_brand'])
g_a_train['brand'] = phone['brand']
g_a_test['brand'] = phone['brand']
Xtr_brand = csr_matrix((np.ones(g_a_train.shape[0]),(g_a_train.trainr, g_a_train.brand)))
Xte_brand = csr_matrix((np.ones(g_a_test.shape[0]),(g_a_test.testr,g_a_test.brand)))
print('Brand features: train shape {}, test shape {}'.format(Xtr_brand.shape, Xte_brand.shape))
```

Brand features: train shape (74645, 131), test shape (112071, 131)

Device model

```
In [70]: phone_model = phone.phone_brand.str.cat(phone.device_model)
model_encoder = LabelEncoder().fit(phone_model)
phone['model'] = model_encoder.transform(phone_model)
g_a_train['model'] = phone['model']
g_a_test['model'] = phone['model']
Xtr_model = csr_matrix((np.ones(g_a_train.shape[0]),
                             (g_a_train.trainr, g_a_train.model)))
Xte_model = csr_matrix((np.ones(g_a_test.shape[0]),
                             (g_a_test.testr, g_a_test.model)))
print('Device model features: train shape {}, test shape {}'.format(Xtr_model.shape, Xte_model.shape))
```

Device model features: train shape (74645, 1667), test shape (112071, 1667)

Installed apps features

For each device we want to have list of installed applications. So we'll have as many feature columns as there are distinct apps.

Apps are linked to devices through events. So we'll do the following:

merge device_id column from events table to app_events group the resulting dataframe by device_id and app and aggregate merge in trainrow and testrow columns to know at which row to put each device in the features matrix

```
In [32]: apps_encoder = LabelEncoder().fit(appevents.app_id)
appevents['app'] = apps_encoder.transform(appevents.app_id)
napps = len(apps_encoder.classes_)
deviceapps = (appevents.merge(events[['device_id']], how='left', left_on='event_id', right_index=True)
               .groupby(['device_id', 'app'])['app'].agg(['size'])
               .merge(g_a_train[['trainr']], how='left', left_index=True, right_index=True)
               .merge(g_a_test[['testr']], how='left', left_index=True, right_index=True)
               .reset_index())
```

In [33]: `deviceapps.head()`

Out[33]:

	device_id	app	size	trainr	testr
0	-9222956879900151005	548	18	21594.0	NaN
1	-9222956879900151005	1096	18	21594.0	NaN
2	-9222956879900151005	1248	26	21594.0	NaN
3	-9222956879900151005	1545	12	21594.0	NaN
4	-9222956879900151005	1664	18	21594.0	NaN

Next step is to build a feature matrix. Data will be all ones, row_ind comes from trainr or testr and col_ind is the label-encoded app_id.

```
In [36]: dfm = deviceapps.dropna(subset=['trainr'])
Xtr_app = csr_matrix((np.ones(dfm.shape[0]), (dfm.trainr, dfm.app)),
                    shape=(g_a_train.shape[0], napps))
dfm = deviceapps.dropna(subset=['testr'])
Xte_app = csr_matrix((np.ones(dfm.shape[0]), (dfm.testr, dfm.app)),
                    shape=(g_a_test.shape[0], napps))
print('Apps data: train shape {}, test shape {}'.format(Xtr_app.shape, Xte_app.shape))
```

Apps data: train shape (74645, 19237), test shape (112071, 19237)

App labels features

We can create app labels merging app_labels with the deviceapps dataframe.

```
In [40]: applabels = applabels.loc[applabels.app_id.isin(appevents.app_id.unique())]
applabels['app'] = apps_encoder.transform(applabels.app_id)
labelencoder = LabelEncoder().fit(applabels.label_id)
applabels['label'] = labelencoder.transform(applabels.label_id)
nlabels = len(labelencoder.classes_)
```

```
In [41]: devicelabels = (deviceapps[['device_id', 'app']]
                    .merge(applabels[['app', 'label']])
                    .groupby(['device_id', 'label'])['app'].agg(['size'])
                    .merge(g_a_train[['trainr']], how='left', left_index=True,
right_index=True)
                    .merge(g_a_test[['testr']], how='left', left_index=True, ri
ght_index=True)
                    .reset_index())
```

In [42]: `devicelabels.head()`

Out[42]:

	device_id	label	size	trainr	testr
0	-9222956879900151005	117	1	21594.0	NaN
1	-9222956879900151005	120	1	21594.0	NaN
2	-9222956879900151005	126	1	21594.0	NaN
3	-9222956879900151005	138	2	21594.0	NaN
4	-9222956879900151005	147	2	21594.0	NaN

```
In [45]: dfm = devicelabels.dropna(subset=['trainr'])
Xtr_label = csr_matrix((np.ones(dfm.shape[0]), (dfm.trainr, dfm.label)),
                        shape=(g_a_train.shape[0],nlabels))
dfm = devicelabels.dropna(subset=['testr'])
Xte_label = csr_matrix((np.ones(dfm.shape[0]), (dfm.testr, dfm.label)),
                        shape=(g_a_test.shape[0],nlabels))
print('Labels data: train shape {}, test shape {}'.format(Xtr_label.shape,
Xte_label.shape))
```

Labels data: train shape (74645, 492), test shape (112071, 492)

Features concatenation

```
In [46]: Xtrain = hstack((Xtr_brand, Xtr_model, Xtr_app, Xtr_label), format='csr')
Xtest = hstack((Xte_brand, Xte_model, Xte_app, Xte_label), format='csr')
print('All features: train shape {}, test shape {}'.format(Xtrain.shape, Xtest.shape))
```

All features: train shape (74645, 21527), test shape (112071, 21527)

Performing cross-validation

```
In [48]: targ_encoder = LabelEncoder().fit(g_a_train.group)
y = targ_encoder.transform(g_a_train.group)
nclasses = len(targ_encoder.classes_)
```




```
In [50]: # Defining loss- score function
def score(clf, random_state = 0):
    kf = StratifiedKFold(y, n_folds=5, shuffle=True, random_state=random_state)
    pred = np.zeros((y.shape[0], nclasses))
    for itrain, itest in kf:
        Xtr, Xte = Xtrain[itrain, :], Xtrain[itest, :]
        ytr, yte = y[itrain], y[itest]
        clf.fit(Xtr, ytr)
        pred[itest, :] = clf.predict_proba(Xte)

    # Resize to one fold (for kernels)
    return log_loss(yte, pred[itest, :])
    print("{:.5f}".format(log_loss(yte, pred[itest, :])), end=' ')
print('')
return log_loss(y, pred)
```

We've tested values for regularization constant C. Since there is probably a lot of columns which are not so important (rare apps or models of brands) we are probably going to get better score with stronger regularization which means that C value will probably going to be below 1.

```
In [61]: cvalue = np.logspace(-3,0,5)
res = []
for C in cvalue:
    res.append(score(LogisticRegression(C = C)))
plt.semilogx(cvalue, res, '-o');
```



So it looks like the best value for C could be between 0.01 and 0.1.

```
In [63]: score(LogisticRegression(C=0.01))
```

```
Out[63]: 2.2848755470140127
```

```
In [55]: score(LogisticRegression(C=0.02))
```

```
Out[55]: 2.2797068236722908
```

```
In [64]: score(LogisticRegression(C=0.03))
```

```
Out[64]: 2.2796060828323981
```

```
In [65]: score(LogisticRegression(C=0.04))
```

```
Out[65]: 2.2809556715503021
```

```
In [66]: score(LogisticRegression(C=0.05))
```

```
Out[66]: 2.2828903616369471
```

LogisticRegression classifier solves multiclass classification problem -in form of one versus rest fashion. But we can also fit a multinomial model that optimizes the multiclass logloss like in our case. We could improve results using this scenario since this is our exact setup.

```
In [62]: score(LogisticRegression(C=0.02, multi_class='multinomial',solver='saga'))
```

```
/root/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/sag.py:32
6: ConvergenceWarning: The max_iter was reached which means the coef_ did n
ot converge
    "the coef_ did not converge", ConvergenceWarning)
```

```
Out[62]: 2.2733450166849916
```

```
In [67]: score(LogisticRegression(C=0.02, multi_class='multinomial',solver='lbfgs'))
```

```
Out[67]: 2.273326572493398
```

```
In [68]: score(LogisticRegression(C=0.02, multi_class='multinomial',solver='newton-c
g'))
```

```
Out[68]: 2.2731559680466482
```

```
In [69]: score(LogisticRegression(C=0.02, multi_class='multinomial',solver='sag'))
```

```
/root/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/sag.py:32
6: ConvergenceWarning: The max_iter was reached which means the coef_ did n
ot converge
    "the coef_ did not converge", ConvergenceWarning)
```

```
Out[69]: 2.273158162504858
```

Test data predictions

```
In [57]: clf = LogisticRegression(C=0.02, multi_class='multinomial',solver='lbfgs')
clf.fit(Xtrain, y)
pred = pd.DataFrame(clf.predict_proba(Xtest), index = g_a_test.index, colum
ns=targ_encoder.classes_)
```

```
In [58]: pred.head()
```

```
Out[58]:
```

	F23-	F24-26	F27-28	F29-32	F33-42	F43+	M23+
device_id							
1002079943728939269	0.001424	0.005998	0.013605	0.013286	0.025313	0.046103	0.001424
-1547860181818787117	0.007414	0.013299	0.031228	0.058677	0.072686	0.151391	0.007414
7374582448058474277	0.023158	0.036713	0.036233	0.158343	0.162774	0.079852	0.023158
-6220210354783429585	0.003474	0.030860	0.008801	0.012351	0.050697	0.172943	0.003474
-5893464122623104785	0.046952	0.065640	0.042578	0.062522	0.056329	0.043467	0.046952

Storing best predictions in CSV file

```
In [59]: pred.to_csv('predictions.csv',index=True)
```