

Data analysis and visualization

```
In [1]: %matplotlib inline
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import os
from mpl_toolkits.basemap import Basemap
```

```
In [2]: # Input data from "/data" director
os.listdir("data")
```

```
Out[2]: ['app_events.csv',
'phone_brand_device_model.csv',
'app_events.csv.zip',
'app_labels.csv.zip',
'app_labels.csv',
'gender_age_test.csv.zip',
'phone_brand_device_model.csv.zip',
'gender_age_test.csv',
'events.csv',
'events.csv.zip',
'sample_submission.csv.zip',
'gender_age_train.csv.zip',
'gender_age_train.csv',
'label_categories.csv.zip',
'sample_submission.csv',
'label_categories.csv']
```

```
In [3]: import seaborn as sns
sns.set(color_codes=True)
app_event=pd.read_csv("data/events.csv")
app_event.shape
```


```
Out[3]: (3252950, 5)
```

```
In [4]: app_event.timestamp=pd.to_datetime(app_event.timestamp)
app_event['time_hour'] = app_event.timestamp.apply(lambda x: x.hour)
```


```
In [5]: # Show frequency of events by hour  
app_event['time_hour'].value_counts()
```

```
Out[5]: 10    183839  
        21    181175  
        20    178179  
        22    168246  
        11    167025  
        19    166160  
         9    166061  
        12    162745  
         8    157896  
        18    156209  
        13    155337  
        17    153516  
        14    151379  
        16    150732  
        15    149912  
         7    146667  
        23    136339  
         0    120512  
         6    116370  
         1     72671  
         5     66411  
         2     53764  
         3     47048  
         4     44757  
Name: time_hour, dtype: int64
```

```
In [6]: ax = sns.countplot(x="time_hour", data=app_event)
```




```
In [7]: import calendar
app_event['week_day'] = app_event.timestamp.apply(lambda x: calendar.day_name[x.weekday()])
ax = sns.countplot(x="week_day", data=app_event)
```




```
In [8]: gender=pd.read_csv("data/gender_age_train.csv")
print(gender.gender.value_counts())
ax = sns.countplot(x="gender", data=gender)
```


```
M    47904
F    26741
Name: gender, dtype: int64
```



```
In [9]: # Distribution by age  
sns.distplot(gender.age, hist=False);
```




```
In [10]: # Distribution by sex  
sns.kdeplot(gender.age[gender.gender=="M"], label="Male")  
sns.kdeplot(gender.age[gender.gender=="F"], label="Female")  
plt.legend();
```



Female at old age are using mobiles little bit more than males at old age


```
In [11]: print("Male mobile usage count by age")
ax = sns.countplot(x="group", data=gender[gender.gender=="M"])
```

Male mobile usage count by age



```
In [12]: print("Female mobile usage count by age")
ax = sns.countplot(x="group", data=gender[gender.gender=="F"])
```

Female mobile usage count by age




```
In [13]: appscategories=pd.read_csv("data/label_categories.csv")
print(appscategories.head())
print(appscategories.shape)
```

```
   label_id  category
0         1         NaN
1         2  game-game type
2         3  game-Game themes
3         4  game-Art Style
4         5  game-Leisure time
(930, 2)
```

Joint visualisation - male and female

```
In [14]: gender.group.value_counts().sort_index(ascending=False).plot(kind='barh')
```


```
Out[14]: <matplotlib.axes._subplots.AxesSubplot at 0x7f31e9ea9a58>
```



Further break down by age

```
In [15]: c = gender.groupby(['age', 'gender']).size().unstack().reindex(index=np.arange(gender.age.min(), gender.age.max()+1)).fillna(0)
ax1, ax2 = c.plot(kind='bar', figsize=(22,6), subplots=True);
ax1.vlines(np.array([23,26,28,32,42])-0.5,0,1800,alpha=0.5,linewidth=1,color='r')
ax2.vlines(np.array([22,26,28,31,38])-0.5,0,3000,alpha=0.5,linewidth=1,color='r')
```

```
Out[15]: <matplotlib.collections.LineCollection at 0x7f31e9dfa438>
```



Locations visualization

```
In [16]: df_events = pd.read_csv("data/events.csv", dtype={'device_id': np.str})
df_events.head()
```

Out[16]:

	event_id	device_id	timestamp	longitude	latitude
0	1	29182687948017175	2016-05-01 00:55:25	121.38	31.24
1	2	-6401643145415154744	2016-05-01 00:54:12	103.65	30.97
2	3	-4833982096941402721	2016-05-01 00:08:05	106.60	29.70
3	4	-6815121365017318426	2016-05-01 00:06:40	104.27	23.28
4	5	-5373797595892518570	2016-05-01 00:07:18	115.88	28.66

Locations of events - World map

Expand size of the screen so larger map can nicely fit without truncated window with scroller.

```
In [17]: %%html
<style>
.output_wrapper, .output {
    height:auto !important;
    max-height:1000px; /* your desired max-height here */
}
.output_scroll {
    box-shadow:none !important;
    webkit-box-shadow:none !important;
}
</style>
```

```
In [18]: # Set plot
df_events_sample = df_events.sample(n=90000)
plt.figure(1, figsize=(20,10))
pd.set_option('display.max_colwidth', -1)

# Map of World
map = Basemap(projection='merc',
              llcrnrlat=-60,
              urcrnrlat=65,
              llcrnrlon=-120,
              urcrnrlon=180,
              lat_ts=0,
              resolution='c')

map.fillcontinents(color='#500000',lake_color='#000000') # grey land, black
Lakes
map.drawmapboundary(fill_color='#202020')                # black background
map.drawcountries(linewidth=0.1, color="w")               # white line of co
untry borders

# Plot the data
mxy = map(df_events_sample["longitude"].tolist(), df_events_sample["latitud
e"].tolist())
map.scatter(mxy[0], mxy[1], s=3, c="#12AABB", lw=0, alpha=1, zorder=5)

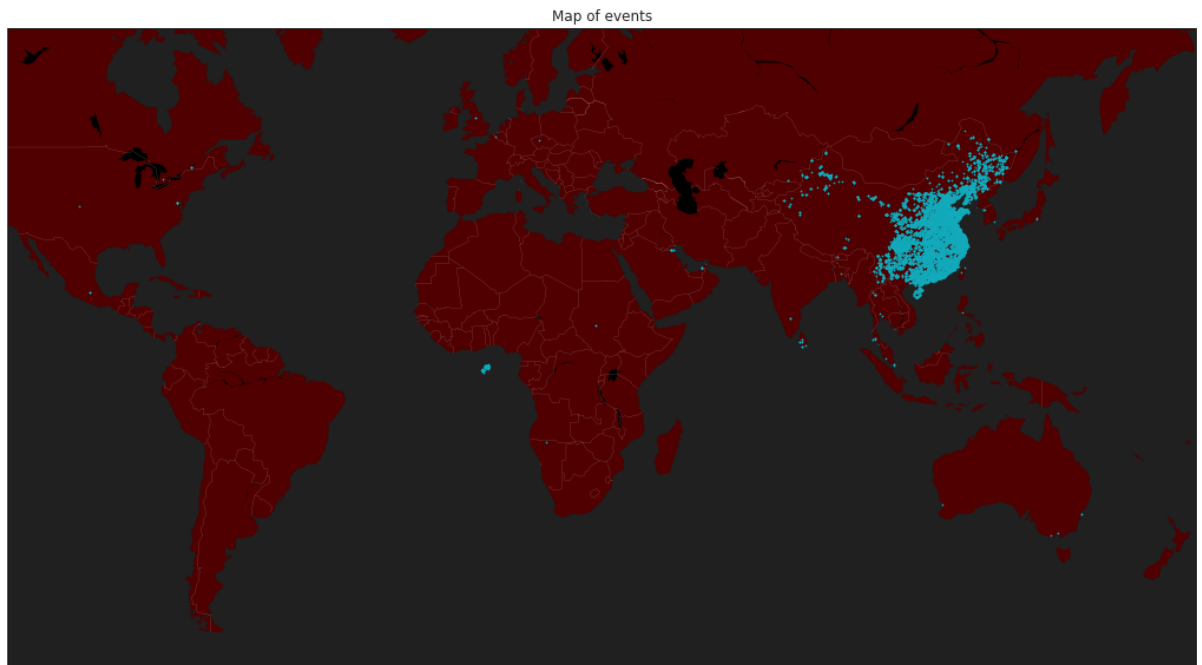
plt.title("Map of events")
plt.show()
```



```

/home/ec2-user/anaconda3/lib/python3.6/site-packages/mpl_toolkits/basemap/_
_init__.py:1767: MatplotlibDeprecationWarning: The get_axis_bgcolor functio
n was deprecated in version 2.0. Use get_facecolor instead.
    axisbgc = ax.get_axis_bgcolor()
/home/ec2-user/anaconda3/lib/python3.6/site-packages/mpl_toolkits/basemap/_
_init__.py:1698: MatplotlibDeprecationWarning: The axesPatch function was d
eprecated in version 2.1. Use Axes.patch instead.
    limb = ax.axesPatch
/home/ec2-user/anaconda3/lib/python3.6/site-packages/mpl_toolkits/basemap/_
_init__.py:3222: MatplotlibDeprecationWarning: The ishold function was depr
ecated in version 2.0.
    b = ax.ishold()
/home/ec2-user/anaconda3/lib/python3.6/site-packages/mpl_toolkits/basemap/_
_init__.py:3231: MatplotlibDeprecationWarning: axes.hold is deprecated.
    See the API Changes document (http://matplotlib.org/api/api\_changes.htm
l)
    for more details.
    ax.hold(b)

```



As we can see some events have $(lat,lon) = (0,0)$ which probably means that location couldn't be determined. We can find all event on that location and all events that have longitud and latitud less than 1 which means they are really close to that location. That is location is not probable since it's points to sea area close to African cost.

```
In [19]: df_at0 = df_events[(df_events["longitude"]==0) & (df_events["latitude"]==0
)]
df_near0 = df_events[(df_events["longitude"]>-1) &\
                      (df_events["longitude"]<1) &\
                      (df_events["latitude"]>-1) &\
                      (df_events["latitude"]<1)]

print("Total number of events:", len(df_events))
print("Number of events at (0,0):", len(df_at0))
print("Number of events near (0,0):", len(df_near0))
```

```
Total number of events: 3252950
Number of events at (0,0): 968675
Number of events near (0,0): 969871
```

Locations of events - zooming in region of China

```
In [20]: # Locate China region
lon_min, lon_max = 75, 135
lat_min, lat_max = 15, 55

idx_china = (df_events["longitude"]>lon_min) &\
             (df_events["longitude"]<lon_max) &\
             (df_events["latitude"]>lat_min) &\
             (df_events["latitude"]<lat_max)

df_events_china = df_events[idx_china].sample(n=100000)

# China
plt.figure(2, figsize=(20,15))

map_zoom = Basemap(projection='merc',
                    llcrnrlat=lat_min,
                    urcrnrlat=lat_max,
                    llcrnrlon=lon_min,
                    urcrnrlon=lon_max,
                    lat_ts=35,
                    resolution='c')

map_zoom.fillcontinents(color='#500000',lake_color='#000000') # dark grey land, black lakes
map_zoom.drawmapboundary(fill_color='#202020')                # black background
map_zoom.drawcountries(linewidth=0.1, color="w")              # thin white line for country borders

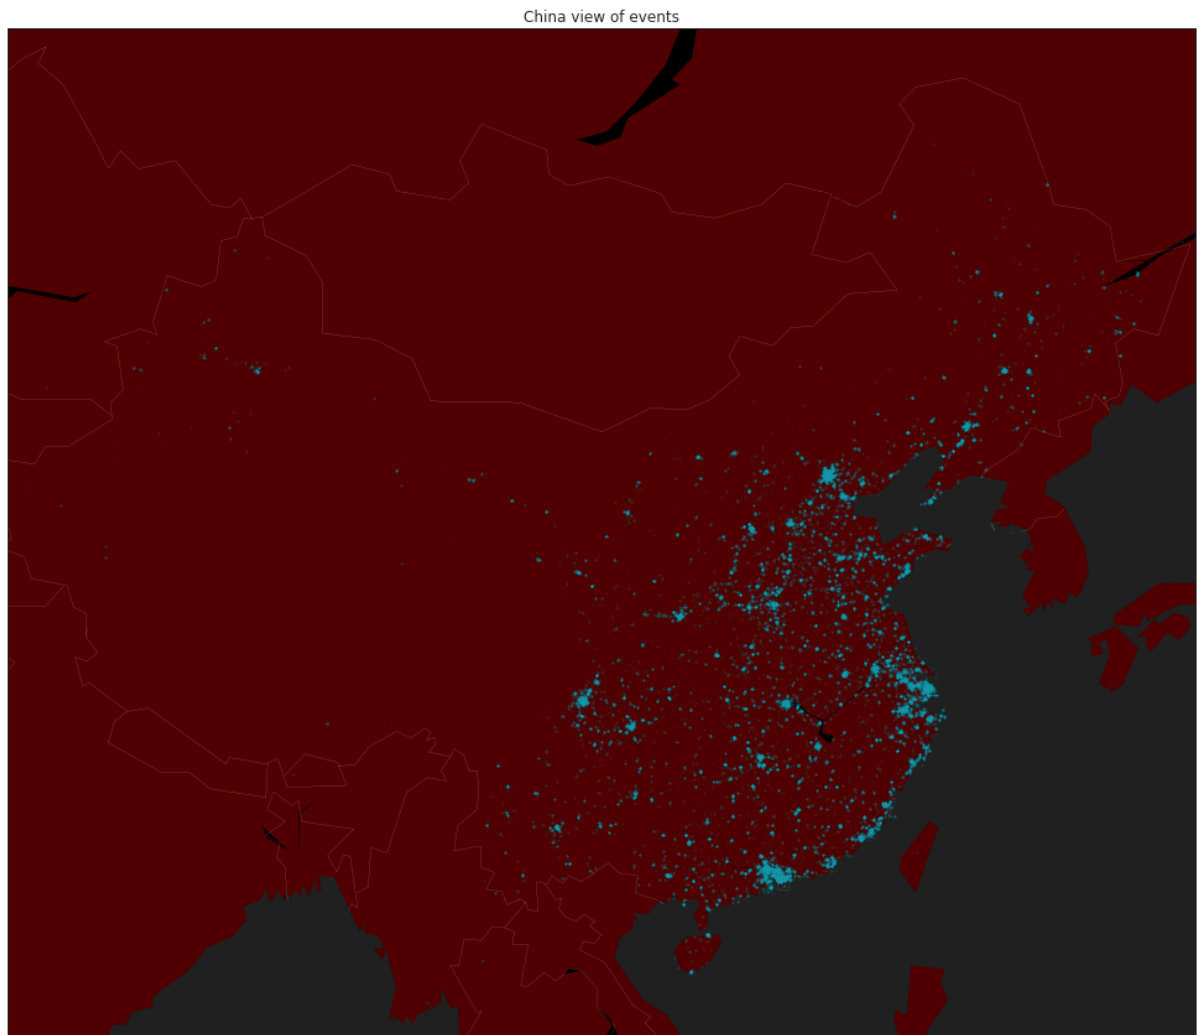
# Plot the data
mxy = map_zoom(df_events_china["longitude"].tolist(), df_events_china["latitude"].tolist())
map_zoom.scatter(mxy[0], mxy[1], s=5, c="#12AABB", lw=0, alpha=0.05, zorder=5)

plt.title("China view of events")
plt.show()
```

```

/home/ec2-user/anaconda3/lib/python3.6/site-packages/mpl_toolkits/basemap/_
_init__.py:1767: MatplotlibDeprecationWarning: The get_axis_bgcolor functio
n was deprecated in version 2.0. Use get_facecolor instead.
    axisbgc = ax.get_axis_bgcolor()
/home/ec2-user/anaconda3/lib/python3.6/site-packages/mpl_toolkits/basemap/_
_init__.py:1698: MatplotlibDeprecationWarning: The axesPatch function was d
eprecated in version 2.1. Use Axes.patch instead.
    limb = ax.axesPatch
/home/ec2-user/anaconda3/lib/python3.6/site-packages/mpl_toolkits/basemap/_
_init__.py:3222: MatplotlibDeprecationWarning: The ishold function was depr
ecated in version 2.0.
    b = ax.ishold()
/home/ec2-user/anaconda3/lib/python3.6/site-packages/mpl_toolkits/basemap/_
_init__.py:3231: MatplotlibDeprecationWarning: axes.hold is deprecated.
    See the API Changes document (http://matplotlib.org/api/api\_changes.htm
l)
    for more details.
    ax.hold(b)

```



This map nicely shows population density of China. We'll analyze one city region in this case we'll take Shanghai. What follows are maps that are showing longitude and latitude areas. In same way, we can zoom on any area on earth for showing more details in that area.

Locations of applications events - region of Shanghai

```
In [21]: # Shanghai sits on 31.2304° N, 121.4737° E
# Sampling wider Shanghai region
lon_min, lon_max = 115, 125
lat_min, lat_max = 28, 35

idx_shanghai = (df_events["longitude"]>lon_min) &\
                (df_events["longitude"]<lon_max) &\
                (df_events["latitude"]>lat_min) &\
                (df_events["latitude"]<lat_max)

df_events_shanghai = df_events[idx_shanghai]

# Map of Shanghai region
plt.figure(3, figsize=(20,15))

m_shanghai = Basemap(projection='merc',
                      llcrnrlat=lat_min,
                      urcrnrlat=lat_max,
                      llcrnrlon=lon_min,
                      urcrnrlon=lon_max,
                      lat_ts=35,
                      resolution='c')

m_shanghai.fillcontinents(color='#500000',lake_color='#000000') # dark land, black lakes
m_shanghai.drawmapboundary(fill_color='#000000') # black background
m_shanghai.drawcountries(linewidth=0.1, color="w") # white line for country borders

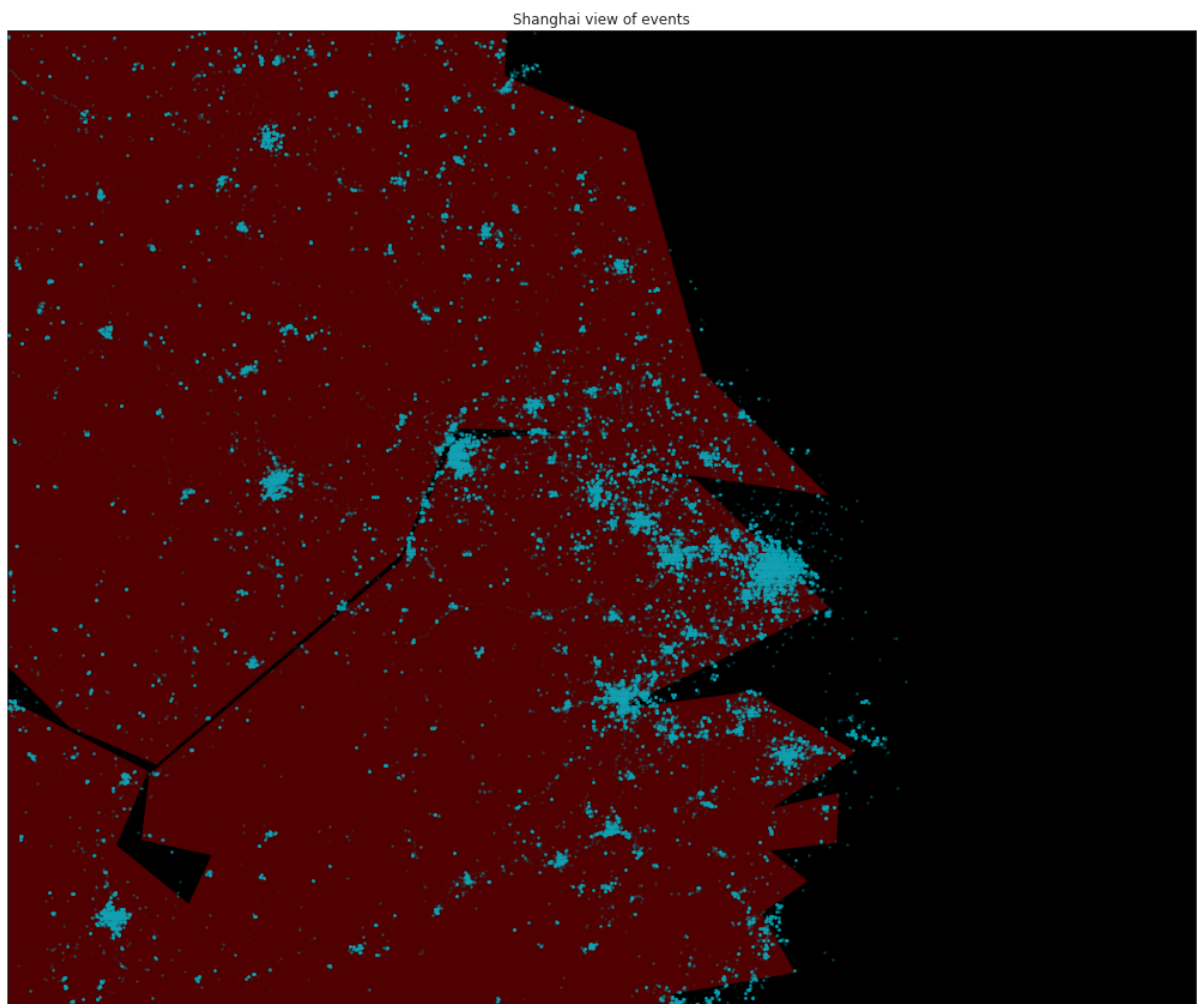
# Plot the data
mxy = m_shanghai(df_events_shanghai["longitude"].tolist(), df_events_shanghai["latitude"].tolist())
m_shanghai.scatter(mxy[0], mxy[1], s=5, c="#12AABB", lw=0, alpha=0.1, zorder=5)

plt.title("Shanghai view of events")
plt.show()
```

```

/home/ec2-user/anaconda3/lib/python3.6/site-packages/mpl_toolkits/basemap/_
_init__.py:1767: MatplotlibDeprecationWarning: The get_axis_bgcolor functio
n was deprecated in version 2.0. Use get_facecolor instead.
    axisbgc = ax.get_axis_bgcolor()
/home/ec2-user/anaconda3/lib/python3.6/site-packages/mpl_toolkits/basemap/_
_init__.py:1698: MatplotlibDeprecationWarning: The axesPatch function was d
eprecated in version 2.1. Use Axes.patch instead.
    limb = ax.axesPatch
/home/ec2-user/anaconda3/lib/python3.6/site-packages/mpl_toolkits/basemap/_
_init__.py:3222: MatplotlibDeprecationWarning: The ishold function was depr
ecated in version 2.0.
    b = ax.ishold()
/home/ec2-user/anaconda3/lib/python3.6/site-packages/mpl_toolkits/basemap/_
_init__.py:3231: MatplotlibDeprecationWarning: axes.hold is deprecated.
    See the API Changes document (http://matplotlib.org/api/api\_changes.htm
1)
    for more details.
    ax.hold(b)

```



We can see that population around big cities is very dispersed.

Now we'll show male and female app events

Male and female app events in region of Shanghai

```
In [22]: # Load the train data and join on the events
df_train = pd.read_csv("data/gender_age_train.csv", dtype={'device_id': np.str})

df_plot = pd.merge(df_train, df_events_shanghai, on="device_id", how="inner")

df_m = df_plot[df_plot["gender"]=="M"]
df_f = df_plot[df_plot["gender"]=="F"]
```

Visualize male and female events in Shanghai


```
In [23]: # Female and male plot

plt.figure(4, figsize=(20,10))

plt.subplot(121)
m_sh_m = Basemap(projection='merc',
                  llcrnrlat=lat_min,
                  urcrnrlat=lat_max,
                  llcrnrlon=lon_min,
                  urcrnrlon=lon_max,
                  lat_ts=35,
                  resolution='c')
m_sh_m.fillcontinents(color='#500000',lake_color='#000000') # dark grey land, black lakes
m_sh_m.drawmapboundary(fill_color='#000000')                # black background
m_sh_m.drawcountries(linewidth=0.1, color="w")              # thin white line for country borders
mxy = m_sh_m(df_m["longitude"].tolist(), df_m["latitude"].tolist())
m_sh_m.scatter(mxy[0], mxy[1], s=5, c="#1292db", lw=0, alpha=0.1, zorder=5)
plt.title("Male events in Shanghai")

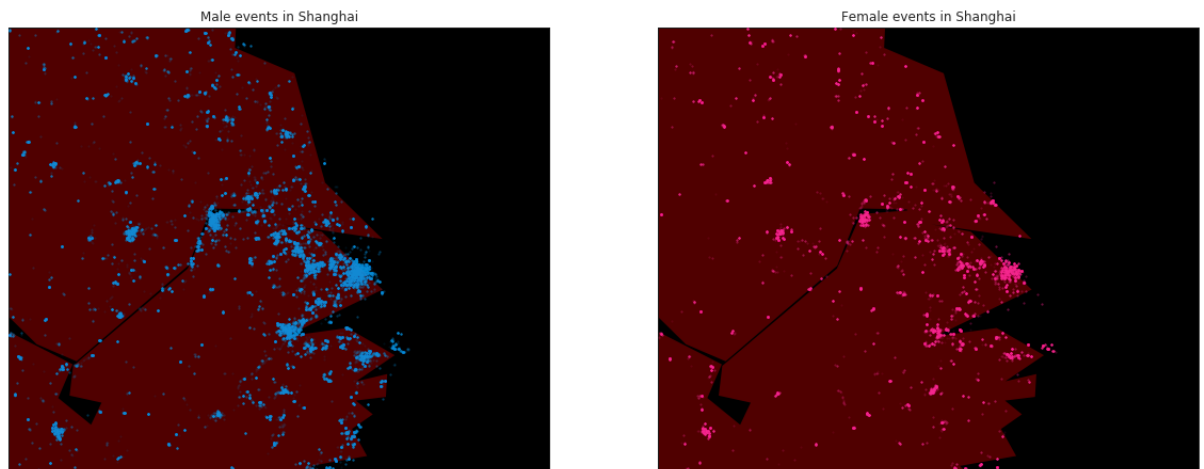
plt.subplot(122)
m_sh_f = Basemap(projection='merc',
                  llcrnrlat=lat_min,
                  urcrnrlat=lat_max,
                  llcrnrlon=lon_min,
                  urcrnrlon=lon_max,
                  lat_ts=35,
                  resolution='c')
m_sh_f.fillcontinents(color='#500000',lake_color='#000000') # dark grey land, black lakes
m_sh_f.drawmapboundary(fill_color='#000000')                # black background
m_sh_f.drawcountries(linewidth=0.1, color="w")              # thin white line for country borders
mxy = m_sh_f(df_f["longitude"].tolist(), df_f["latitude"].tolist())
m_sh_f.scatter(mxy[0], mxy[1], s=5, c="#fd3096", lw=0, alpha=0.1, zorder=5)
plt.title("Female events in Shanghai")

plt.show()
```

```

/home/ec2-user/anaconda3/lib/python3.6/site-packages/mpl_toolkits/basemap/_
_init__.py:1767: MatplotlibDeprecationWarning: The get_axis_bgcolor functio
n was deprecated in version 2.0. Use get_facecolor instead.
    axisbgc = ax.get_axis_bgcolor()
/home/ec2-user/anaconda3/lib/python3.6/site-packages/mpl_toolkits/basemap/_
_init__.py:1698: MatplotlibDeprecationWarning: The axesPatch function was d
eprecated in version 2.1. Use Axes.patch instead.
    limb = ax.axesPatch
/home/ec2-user/anaconda3/lib/python3.6/site-packages/mpl_toolkits/basemap/_
_init__.py:3222: MatplotlibDeprecationWarning: The ishold function was depr
ecated in version 2.0.
    b = ax.ishold()
/home/ec2-user/anaconda3/lib/python3.6/site-packages/mpl_toolkits/basemap/_
_init__.py:3231: MatplotlibDeprecationWarning: axes.hold is deprecated.
    See the API Changes document (http://matplotlib.org/api/api\_changes.htm
l)
    for more details.
    ax.hold(b)

```



For marketing analysis, this might be interesting for further exploration. Which city areas are showing more men activities and which are showing more female activities and in which times of day?

Analysis

Problem classification

Our task is to build a model predicting users' demographic characteristics based on their app usage, geolocation, and mobile device properties. So we need to solve multiclass classification problem. This is a case where one label needs to be predicted based on several others.

Logistic regression

Logistic regression algorithm could be obvious choice for that. In the multiclass case, the training algorithm uses the one-vs-rest (OvR) scheme.

```
In [24]: from sklearn.preprocessing import LabelEncoder
from scipy.sparse import csr_matrix, hstack
from sklearn.linear_model import LogisticRegression
from sklearn.cross_validation import StratifiedKFold
from sklearn.metrics import log_loss
```

```
/home/ec2-user/anaconda3/lib/python3.6/site-packages/sklearn/cross_validation.py:41: DeprecationWarning: This module was deprecated in version 0.18 in favor of the model_selection module into which all the refactored classes and functions are moved. Also note that the interface of the new CV iterators are different from that of this module. This module will be removed in 0.20.
```

```
"This module will be removed in 0.20.", DeprecationWarning)
```

Loading of data

```
In [25]: # Read gender Train and Test data
datadir = 'data'
g_a_train = pd.read_csv(os.path.join(datadir, 'gender_age_train.csv'), index_col='device_id')
g_a_test = pd.read_csv(os.path.join(datadir, 'gender_age_test.csv'), index_col='device_id')
phone = pd.read_csv(os.path.join(datadir, 'phone_brand_device_model.csv'))

print(g_a_train.head())
print("-----")
print(g_a_test.head())
print("-----")
print(phone.head())
```

```

              gender  age  group
device_id
-8076087639492063270  M    35  M32-38
-2897161552818060146  M    35  M32-38
-8260683887967679142  M    35  M32-38
-4938849341048082022  M    30  M29-31
 245133531816851882  M    30  M29-31
-----
Empty DataFrame
Columns: []
Index: [1002079943728939269, -1547860181818787117, 7374582448058474277, -6220210354783429585, -5893464122623104785]
-----
              device_id  phone_brand  device_model
0 -8890648629457979026   小米         红米
1  1277779817574759137   小米         MI 2
2  5137427614288105724   三星        Galaxy S4
3  3669464369358936369  SUGAR      时尚手机
4 -5019277647504317457   三星        Galaxy Note 2
```

```
In [26]: # Remove duplicate device ids in the phones
phone = phone.drop_duplicates(subset=['device_id'], keep='first').set_index('device_id')
events = pd.read_csv('data/events.csv', parse_dates=['timestamp'], index_col='event_id')
appevents = pd.read_csv('data/app_events.csv', usecols=['event_id', 'app_id', 'is_active'], dtype={'is_active': bool})
applabels = pd.read_csv('data/app_labels.csv')
```

```
/home/ec2-user/anaconda3/lib/python3.6/site-packages/numpy/lib/arraysetops.py:463: FutureWarning: elementwise comparison failed; returning scalar instead, but in the future will perform elementwise comparison
  mask |= (ar1 == a)
```

```
In [27]: print(phone.head())
print("-----")
print(events.head())
print("-----")
print(appevents.head())
print("-----")
print(applabels.head())
print("-----")
```

	phone_brand	device_model
device_id		
-8890648629457979026	小米	红米
1277779817574759137	小米	MI 2
5137427614288105724	三星	Galaxy S4
3669464369358936369	SUGAR	时尚手机
-5019277647504317457	三星	Galaxy Note 2

	device_id	timestamp	longitude	latitude
event_id				
1	29182687948017175	2016-05-01 00:55:25	121.38	31.24
2	-6401643145415154744	2016-05-01 00:54:12	103.65	30.97
3	-4833982096941402721	2016-05-01 00:08:05	106.60	29.70
4	-6815121365017318426	2016-05-01 00:06:40	104.27	23.28
5	-5373797595892518570	2016-05-01 00:07:18	115.88	28.66

	event_id	app_id	is_active
0	2	5927333115845830913	True
1	2	-5720078949152207372	False
2	2	-1633887856876571208	False
3	2	-653184325010919369	True
4	2	8693964245073640147	True

	app_id	label_id
0	7324884708820027918	251
1	-4494216993218550286	251
2	6058196446775239644	406
3	6058196446775239644	407
4	8694625920731541625	406

Main feature selection

Main features chosen are:

- phone brand
- device model
- installed apps
- app labels

We need to one-hot encode everything and put in sparse matrices which will help deal with a very large number of features. Regarding "Phone brand" feature; we'll make two columns that show which train or test set row a particular device_id belongs to.

```
In [28]: g_a_train['trainr'] = np.arange(g_a_train.shape[0])
         g_a_test['testr'] = np.arange(g_a_test.shape[0])
```

```
In [29]: print(g_a_train.head())
         print("-----")
         print(g_a_test.head())
```

device_id	gender	age	group	trainr
-8076087639492063270	M	35	M32-38	0
-2897161552818060146	M	35	M32-38	1
-8260683887967679142	M	35	M32-38	2
-4938849341048082022	M	30	M29-31	3
245133531816851882	M	30	M29-31	4

device_id	testr
1002079943728939269	0
-1547860181818787117	1
7374582448058474277	2
-6220210354783429585	3
-5893464122623104785	4

Constructing sparse matrix of features in following way:

`csr_matrix((data, (row_ind, col_ind)), [shape=(M, N)])` where data, row_ind and col_ind satisfy the relationship `a[row_ind[k], col_ind[k]] = data[k]`

This allows us to define what values to put into certain places in a sparse matrix. For phone brand data the data array will be all ones, row_ind will be the row number of a device and col_ind will be the number of brand.

Brand features

```
In [30]: brand_encoder = LabelEncoder().fit(phone.phone_brand)
phone['brand'] = brand_encoder.transform(phone['phone_brand'])
g_a_train['brand'] = phone['brand']
g_a_test['brand'] = phone['brand']
Xtr_brand = csr_matrix((np.ones(g_a_train.shape[0]),(g_a_train.trainr, g_a_train.brand)))
Xte_brand = csr_matrix((np.ones(g_a_test.shape[0]),(g_a_test.testr,g_a_test.brand)))
print('Brand features: train shape {}, test shape {}'.format(Xtr_brand.shape, Xte_brand.shape))
```

Brand features: train shape (74645, 131), test shape (112071, 131)

Device model

```
In [31]: phone_model = phone.phone_brand.str.cat(phone.device_model)
model_encoder = LabelEncoder().fit(phone_model)
phone['model'] = model_encoder.transform(phone_model)
g_a_train['model'] = phone['model']
g_a_test['model'] = phone['model']
Xtr_model = csr_matrix((np.ones(g_a_train.shape[0]),
                             (g_a_train.trainr, g_a_train.model)))
Xte_model = csr_matrix((np.ones(g_a_test.shape[0]),
                             (g_a_test.testr, g_a_test.model)))
print('Device model features: train shape {}, test shape {}'.format(Xtr_model.shape, Xte_model.shape))
```

Device model features: train shape (74645, 1667), test shape (112071, 1667)

Installed apps features

For each device we want to have list of installed applications. So we'll have as many feature columns as there are distinct apps.

Apps are linked to devices through events. So we'll do the following:

merge device_id column from events table to app_events group the resulting dataframe by device_id and app and aggregate merge in trainrow and testrow columns to know at which row to put each device in the features matrix

```
In [32]: apps_encoder = LabelEncoder().fit(appevents.app_id)
appevents['app'] = apps_encoder.transform(appevents.app_id)
napps = len(apps_encoder.classes_)
deviceapps = (appevents.merge(events[['device_id']], how='left', left_on='event_id', right_index=True)
               .groupby(['device_id', 'app'])['app'].agg(['size'])
               .merge(g_a_train[['trainr']], how='left', left_index=True, right_index=True)
               .merge(g_a_test[['testr']], how='left', left_index=True, right_index=True)
               .reset_index())
```

In [33]: `deviceapps.head()`

Out[33]:

	device_id	app	size	trainr	testr
0	-9222956879900151005	548	18	21594.0	NaN
1	-9222956879900151005	1096	18	21594.0	NaN
2	-9222956879900151005	1248	26	21594.0	NaN
3	-9222956879900151005	1545	12	21594.0	NaN
4	-9222956879900151005	1664	18	21594.0	NaN

Next step is to build a feature matrix. Data will be all ones, row_ind comes from trainr or testr and col_ind is the label-encoded app_id.

```
In [34]: dfm = deviceapps.dropna(subset=['trainr'])
Xtr_app = csr_matrix((np.ones(dfm.shape[0]), (dfm.trainr, dfm.app)),
                    shape=(g_a_train.shape[0], napps))
dfm = deviceapps.dropna(subset=['testr'])
Xte_app = csr_matrix((np.ones(dfm.shape[0]), (dfm.testr, dfm.app)),
                    shape=(g_a_test.shape[0], napps))
print('Apps data: train shape {}, test shape {}'.format(Xtr_app.shape, Xte_app.shape))
```

Apps data: train shape (74645, 19237), test shape (112071, 19237)

App labels features

We can create app labels merging app_labels with the deviceapps dataframe.

```
In [35]: applabels = applabels.loc[applabels.app_id.isin(appevents.app_id.unique())]
applabels['app'] = apps_encoder.transform(applabels.app_id)
labelencoder = LabelEncoder().fit(applabels.label_id)
applabels['label'] = labelencoder.transform(applabels.label_id)
nlabels = len(labelencoder.classes_)
```

```
In [36]: devicelabels = (deviceapps[['device_id', 'app']]
                        .merge(applabels[['app', 'label']])
                        .groupby(['device_id', 'label'])['app'].agg(['size'])
                        .merge(g_a_train[['trainr']], how='left', left_index=True,
right_index=True)
                        .merge(g_a_test[['testr']], how='left', left_index=True, ri
ght_index=True)
                        .reset_index())
```

In [37]: `devicelabels.head()`

Out[37]:

	device_id	label	size	trainr	testr
0	-9222956879900151005	117	1	21594.0	NaN
1	-9222956879900151005	120	1	21594.0	NaN
2	-9222956879900151005	126	1	21594.0	NaN
3	-9222956879900151005	138	2	21594.0	NaN
4	-9222956879900151005	147	2	21594.0	NaN

```
In [38]: dfm = devicelabels.dropna(subset=['trainr'])
Xtr_label = csr_matrix((np.ones(dfm.shape[0]), (dfm.trainr, dfm.label)),
                        shape=(g_a_train.shape[0],nlabels))
dfm = devicelabels.dropna(subset=['testr'])
Xte_label = csr_matrix((np.ones(dfm.shape[0]), (dfm.testr, dfm.label)),
                        shape=(g_a_test.shape[0],nlabels))
print('Labels data: train shape {}, test shape {}'.format(Xtr_label.shape,
Xte_label.shape))
```

Labels data: train shape (74645, 492), test shape (112071, 492)

Features concatenation

```
In [39]: Xtrain = hstack((Xtr_brand, Xtr_model, Xtr_app, Xtr_label), format='csr')
Xtest = hstack((Xte_brand, Xte_model, Xte_app, Xte_label), format='csr')
print('All features: train shape {}, test shape {}'.format(Xtrain.shape, Xtest.shape))
```

All features: train shape (74645, 21527), test shape (112071, 21527)

Performing cross-validation

```
In [40]: targ_encoder = LabelEncoder().fit(g_a_train.group)
y = targ_encoder.transform(g_a_train.group)
nclasses = len(targ_encoder.classes_)
```




```
In [42]: # Defining loss- score function
def score(clf, random_state = 0):
    kf = StratifiedKFold(y, n_folds=5, shuffle=True, random_state=random_state)
    pred = np.zeros((y.shape[0], nclasses))
    for itrain, itest in kf:
        Xtr, Xte = Xtrain[itrain, :], Xtrain[itest, :]
        ytr, yte = y[itrain], y[itest]
        clf.fit(Xtr, ytr)
        pred[itest, :] = clf.predict_proba(Xte)

    # Resize to one fold (for kernels)
    return log_loss(yte, pred[itest, :])
    print("{:.5f}".format(log_loss(yte, pred[itest, :])), end=' ')
print('')
return log_loss(y, pred)
```

We've tested values for regularization constant C. Since there is probably a lot of columns which are not so important (rare apps or models of brands) we are probably going to get better score with stronger regularization which means that C value will probably going to be below 1.

```
In [43]: cvalue = np.logspace(-3,0,6)
res = []
for C in cvalue:
    res.append(score(LogisticRegression(C = C)))
plt.semilogx(cvalue, res, '-o');
```



So it looks like the best value for C could be between 0.01 and 0.1.

```
In [44]: score(LogisticRegression(C=0.01))
```

```
Out[44]: 2.2848755470140127
```

```
In [45]: score(LogisticRegression(C=0.02))
```

```
Out[45]: 2.2797068236722908
```

```
In [46]: score(LogisticRegression(C=0.03))
```

```
Out[46]: 2.2796060828323981
```

```
In [47]: score(LogisticRegression(C=0.04))
```

```
Out[47]: 2.2809556715503021
```

```
In [48]: score(LogisticRegression(C=0.05))
```

```
Out[48]: 2.2828903616369471
```

LogisticRegression classifier solves multiclass classification problem -in form of one versus rest fashion. But we can also fit a multinomial model that optimizes the multiclass logloss like in our case. We could improve results using this scenario since this is our exact setup.

```
In [49]: score(LogisticRegression(C=0.02, multi_class='multinomial',solver='saga'))
```

```
/home/ec2-user/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/sag.py:326: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
  "the coef_ did not converge", ConvergenceWarning)
```

```
Out[49]: 2.2733463394827584
```

```
In [50]: score(LogisticRegression(C=0.02, multi_class='multinomial',solver='lbfgs'))
```

```
Out[50]: 2.273326572493398
```

```
In [51]: score(LogisticRegression(C=0.02, multi_class='multinomial',solver='newton-cg'))
```

```
Out[51]: 2.2731559680466482
```

```
In [52]: score(LogisticRegression(C=0.02, multi_class='multinomial',solver='sag'))
```

```
/home/ec2-user/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/sag.py:326: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
  "the coef_ did not converge", ConvergenceWarning)
```

```
Out[52]: 2.2731581798416354
```

Benchmark comparison - Test dataset with XGBoost

```
In [3]: import datetime
import pandas as pd
import numpy as np
from sklearn.cross_validation import train_test_split
import xgboost as xgb
```

```

import random
import zipfile
import time
import shutil
from sklearn.metrics import log_loss

random.seed(2016)

def run_xgb(train, test, features, target, random_state=0):
    eta = 0.1
    max_depth = 3
    subsample = 0.7
    colsample_bytree = 0.7
    start_time = time.time()

    print('XGBoost params. ETA: {}, MAX_DEPTH: {}, SUBSAMPLE: {}, COLSAM
    PLE_BY_TREE: {}'.format(eta, max_depth, subsample, colsample_bytree))
    params = {
        "objective": "multi:softprob",
        "num_class": 12,
        "booster": "gbtree",
        "eval_metric": "mlogloss",
        "eta": eta,
        "max_depth": max_depth,
        "subsample": subsample,
        "colsample_bytree": colsample_bytree,
        "silent": 1,
        "seed": random_state,
    }
    num_boost_round = 500
    early_stopping_rounds = 50
    test_size = 0.3

    X_train, X_valid = train_test_split(train, test_size=test_size, rand
    om_state=random_state)
    print('Length train:', len(X_train.index))
    print('Length valid:', len(X_valid.index))
    y_train = X_train[target]
    y_valid = X_valid[target]
    dtrain = xgb.DMatrix(X_train[features], y_train)
    dvalid = xgb.DMatrix(X_valid[features], y_valid)

    watchlist = [(dtrain, 'train'), (dvalid, 'eval')]
    gbm = xgb.train(params, dtrain, num_boost_round, evals=watchlist, ea
    rly_stopping_rounds=early_stopping_rounds, verbose_eval=True)

    print("Validating...")
    check = gbm.predict(xgb.DMatrix(X_valid[features]), ntree_limit=gbm.
    best_iteration)
    score = log_loss(y_valid.tolist(), check)

    print("Predict test set...")
    test_prediction = gbm.predict(xgb.DMatrix(test[features]), ntree_lim
    it=gbm.best_iteration)

    print('Training time: {} minutes'.format(round((time.time() - start_
    time)/60, 2)))

```

```

    return test_prediction.tolist(), score

def create_submission(score, test, prediction):
    # Make Submission
    now = datetime.datetime.now()
    sub_file = 'submission_' + str(score) + '_' + str(now.strftime("%Y-%m-%d-%H-%M")) + '.csv'
    print('Writing submission: ', sub_file)
    f = open(sub_file, 'w')
    f.write('device_id,F23-,F24-26,F27-28,F29-32,F33-42,F43+,M22-,M23-26,M27-28,M29-31,M32-38,M39+\n')
    total = 0
    test_val = test['device_id'].values
    for i in range(len(test_val)):
        str1 = str(test_val[i])
        for j in range(12):
            str1 += ',' + str(prediction[i][j])
        str1 += '\n'
        total += 1
        f.write(str1)
    f.close()

def map_column(table, f):
    labels = sorted(table[f].unique())
    mappings = dict()
    for i in range(len(labels)):
        mappings[labels[i]] = i
    table = table.replace({f: mappings})
    return table

def read_train_test():
    # Events
    print('Read events data')
    events = pd.read_csv("data/events.csv", dtype={'device_id': np.str})
    events['counts'] = events.groupby(['device_id'])['event_id'].transform('count')
    events_small = events[['device_id', 'counts']].drop_duplicates('device_id', keep='first')

    # Phone brand
    print('Read brands data')
    pbd = pd.read_csv("data/phone_brand_device_model.csv", dtype={'device_id': np.str})
    pbd.drop_duplicates('device_id', keep='first', inplace=True)
    pbd = map_column(pbd, 'phone_brand')
    pbd = map_column(pbd, 'device_model')

    # Train
    print('Read training data')
    train = pd.read_csv("data/gender_age_train.csv", dtype={'device_id': np.str})
    train = map_column(train, 'group')
    train = train.drop(['age'], axis=1)
    train = train.drop(['gender'], axis=1)

```

```
train = pd.merge(train, pbd, how='left', on='device_id', left_index=
True)
train = pd.merge(train, events_small, how='left', on='device_id', le
ft_index=True)
train.fillna(-1, inplace=True)

# Test
print('Read test data')
test = pd.read_csv("data/gender_age_test.csv", dtype={'device_id': n
p.str})
test = pd.merge(test, pbd, how='left', on='device_id', left_index=Tr
ue)
test = pd.merge(test, events_small, how='left', on='device_id', left
_index=True)
test.fillna(-1, inplace=True)

# Features
features = list(test.columns.values)
features.remove('device_id')

return train, test, features

train, test, features = read_train_test()
print('Length of train: ', len(train))
print('Length of test: ', len(test))
print('Features [{ }]: {}'.format(len(features), sorted(features)))
test_prediction, score = run_xgb(train, test, features, 'group')
print("LS: {}".format(round(score, 5)))
create_submission(score, test, test_prediction)
```

```
Read events data
Read brands data
Read training data
Read test data
Length of train: 74645
Length of test: 112071
Features [3]: ['counts', 'device_model', 'phone_brand']
XGBoost params. ETA: 0.1, MAX_DEPTH: 3, SUBSAMPLE: 0.7, COLSAMPLE_BY_TREE: 0.7
Length train: 52251
Length valid: 22394
[0] train-mlogloss:2.47571 eval-mlogloss:2.47616
Multiple eval metrics have been passed: 'eval-mlogloss' will be used for
early stopping.
```

Will train until eval-mlogloss hasn't improved in 50 rounds.

```
[1] train-mlogloss:2.46755 eval-mlogloss:2.46858
[2] train-mlogloss:2.46021 eval-mlogloss:2.46177
[3] train-mlogloss:2.45349 eval-mlogloss:2.45551
[4] train-mlogloss:2.4475 eval-mlogloss:2.44996
[5] train-mlogloss:2.4422 eval-mlogloss:2.44512
[6] train-mlogloss:2.43738 eval-mlogloss:2.44074
[7] train-mlogloss:2.43297 eval-mlogloss:2.43675
[8] train-mlogloss:2.42895 eval-mlogloss:2.43312
[9] train-mlogloss:2.42525 eval-mlogloss:2.42983
[10] train-mlogloss:2.42203 eval-mlogloss:2.42711
[11] train-mlogloss:2.41906 eval-mlogloss:2.4245
[12] train-mlogloss:2.41634 eval-mlogloss:2.42212
[13] train-mlogloss:2.41398 eval-mlogloss:2.42011
[14] train-mlogloss:2.41178 eval-mlogloss:2.41839
[15] train-mlogloss:2.40971 eval-mlogloss:2.41665
[16] train-mlogloss:2.40777 eval-mlogloss:2.41508
[17] train-mlogloss:2.40596 eval-mlogloss:2.41359
[18] train-mlogloss:2.40433 eval-mlogloss:2.41223
[19] train-mlogloss:2.40279 eval-mlogloss:2.41103
[20] train-mlogloss:2.40146 eval-mlogloss:2.41
[21] train-mlogloss:2.40015 eval-mlogloss:2.40903
[22] train-mlogloss:2.3989 eval-mlogloss:2.4081
[23] train-mlogloss:2.39779 eval-mlogloss:2.40738
[24] train-mlogloss:2.39677 eval-mlogloss:2.40656
[25] train-mlogloss:2.39576 eval-mlogloss:2.40586
[26] train-mlogloss:2.39496 eval-mlogloss:2.40534
[27] train-mlogloss:2.39421 eval-mlogloss:2.40485
[28] train-mlogloss:2.39336 eval-mlogloss:2.40437
[29] train-mlogloss:2.39266 eval-mlogloss:2.40399
[30] train-mlogloss:2.39194 eval-mlogloss:2.40356
[31] train-mlogloss:2.39125 eval-mlogloss:2.40323
[32] train-mlogloss:2.39069 eval-mlogloss:2.40295
[33] train-mlogloss:2.39008 eval-mlogloss:2.40264
[34] train-mlogloss:2.38951 eval-mlogloss:2.40238
[35] train-mlogloss:2.38893 eval-mlogloss:2.40205
[36] train-mlogloss:2.38842 eval-mlogloss:2.40175
[37] train-mlogloss:2.38799 eval-mlogloss:2.40155
[38] train-mlogloss:2.38752 eval-mlogloss:2.4013
[39] train-mlogloss:2.38705 eval-mlogloss:2.40119
[40] train-mlogloss:2.38662 eval-mlogloss:2.40101
[41] train-mlogloss:2.38616 eval-mlogloss:2.4008
```

Test data predictions - LogisticRegression

```
In [52]: clf = LogisticRegression(C=0.02, multi_class='multinomial', solver='lbfgs')
         clf.fit(X_train, y_train)
         pred = pd.DataFrame(clf.predict_proba(X_test), index=g_a_test.index, columns=targ_encoder.classes_)
         print('train-mlogloss:', train_mlogloss, 'eval-mlogloss:', eval_mlogloss)
```

```
In [53]: print('train-mlogloss:', train_mlogloss, 'eval-mlogloss:', eval_mlogloss)
```

```
Out[53]:
```

	F23	F24-26	F27-28	F29-32	F33-42	F43+	M23
device_id							
1002079943728939269	0.001424	0.005998	0.013605	0.013286	0.025313	0.046103	0.0
154786018181878717	0.007496	0.013299	0.031228	0.058677	0.072686	0.151391	0.0
7374582448058474277	0.023158	0.036713	0.036233	0.158343	0.162774	0.079852	0.0
6220210354783429585	0.003474	0.030860	0.068801	0.012351	0.050697	0.172943	0.0
58934641122623104785	0.046982	0.065640	0.042578	0.067222	0.056329	0.043467	0.0

Storing best predictions in CSV file

```
In [54]: pred.to_csv('predictions.csv', index=True)
```

Free form data visualizations


```

[99] train-mlogloss:2.37127 eval-mlogloss:2.39769
[100] train-mlogloss:2.37106 eval-mlogloss:2.39767
[101] train-mlogloss:2.37085 eval-mlogloss:2.39762
[102] train-mlogloss:2.37061 eval-mlogloss:2.39756
[103] train-mlogloss:2.37044 eval-mlogloss:2.39752
[104] train-mlogloss:2.37023 eval-mlogloss:2.39745
[105] train-mlogloss:2.37003 eval-mlogloss:2.39748
[106] train-mlogloss:2.36982 eval-mlogloss:2.39747
[107] train-mlogloss:2.36964 eval-mlogloss:2.39754
[108] train-mlogloss:2.36945 eval-mlogloss:2.39757
[109] train-mlogloss:2.36928 eval-mlogloss:2.3976
[110] train-mlogloss:2.36908 eval-mlogloss:2.39755
[111] train-mlogloss:2.36892 eval-mlogloss:2.3976
[112] train-mlogloss:2.3687 eval-mlogloss:2.3976
[113] train-mlogloss:2.36848 eval-mlogloss:2.39758
[114] train-mlogloss:2.36834 eval-mlogloss:2.39757
[115] train-mlogloss:2.36818 eval-mlogloss:2.39755
[116] train-mlogloss:2.36802 eval-mlogloss:2.39754
[117] train-mlogloss:2.36782 eval-mlogloss:2.3975
[118] train-mlogloss:2.36767 eval-mlogloss:2.39746
[119] train-mlogloss:2.36749 eval-mlogloss:2.39745
[120] train-mlogloss:2.36731 eval-mlogloss:2.39745
[121] train-mlogloss:2.36714 eval-mlogloss:2.39746
[122] train-mlogloss:2.36695 eval-mlogloss:2.39748
[123] train-mlogloss:2.36675 eval-mlogloss:2.39747
[124] train-mlogloss:2.36655 eval-mlogloss:2.39744
[125] train-mlogloss:2.36637 eval-mlogloss:2.39743
[126] train-mlogloss:2.36616 eval-mlogloss:2.39735
[127] train-mlogloss:2.36599 eval-mlogloss:2.39735
[128] train-mlogloss:2.3658 eval-mlogloss:2.39723
[129] train-mlogloss:2.36561 eval-mlogloss:2.39722
[130] train-mlogloss:2.36546 eval-mlogloss:2.39725
[131] train-mlogloss:2.3653 eval-mlogloss:2.39726
[132] train-mlogloss:2.36514 eval-mlogloss:2.39721
[133] train-mlogloss:2.36495 eval-mlogloss:2.39721
[134] train-mlogloss:2.36482 eval-mlogloss:2.3972
[135] train-mlogloss:2.36462 eval-mlogloss:2.3972
[136] train-mlogloss:2.36445 eval-mlogloss:2.39715
[137] train-mlogloss:2.36428 eval-mlogloss:2.39711

```

In [4]:

```

[138] author train-mlogloss:2.36414 eval-mlogloss:2.3971
[139] author train-mlogloss:2.36397 eval-mlogloss:2.39709
[140] author train-mlogloss:2.36381 eval-mlogloss:2.39709
[141] # Z-Turbo def the first 3 features eval-mlogloss:2.39706
[142] # I added two new features and t-SNE clustering & visualization eval-mlogloss:2.39703
[143] # used some ideas from https://www.kaggle.com/cast42/santander-customer- eval-mlogloss:2.39701
[144] satisfaction/t-sne-manifold-visualisation/code eval-mlogloss:2.39705
[145] train-mlogloss:2.36295 eval-mlogloss:2.39705
[146] import pandas as pd train-mlogloss:2.36276 eval-mlogloss:2.39704
[147] import numpy as np train-mlogloss:2.36259 eval-mlogloss:2.39703
[148] from sklearn.cross_validation import train_test_split eval-mlogloss:2.39702
[149] from sklearn import manifold train-mlogloss:2.36227 eval-mlogloss:2.39705
[150] from sklearn.cross_validation import StratifiedShuffleSplit eval-mlogloss:2.39707
[151] from sklearn.preprocessing import normalize eval-mlogloss:2.39708
[152] import matplotlib.pyplot as plt train-mlogloss:2.36197 eval-mlogloss:2.39709
[153] import matplotlib.cm as cm train-mlogloss:2.36181 eval-mlogloss:2.39701
[154] train-mlogloss:2.36144 eval-mlogloss:2.39696
[155] train-mlogloss:2.36125 eval-mlogloss:2.39694

```

```

156] train_mlogloss:2.39695 eval_mlogloss:2.39695
157] ss_train_mlogloss:2.39694 eval_mlogloss:2.39694
158] for train_mlogloss:2.39694 eval_mlogloss:2.39694
159] break_mlogloss:2.36066 eval_mlogloss:2.39695
160] train_mlogloss:2.36048 eval_mlogloss:2.39693
161] X_train_mlogloss:2.36048 eval_mlogloss:2.39693
162] Y_train_mlogloss:2.36048 eval_mlogloss:2.39693
163] train_mlogloss:2.36004 eval_mlogloss:2.39697
164] train_mlogloss:2.36004 eval_mlogloss:2.39692
165] train_mlogloss:2.36004 eval_mlogloss:2.39686
166] train_mlogloss:2.35956 eval_mlogloss:2.39684
167] train_mlogloss:2.35956 eval_mlogloss:2.39677
168] train_mlogloss:2.35956 eval_mlogloss:2.39677
169] train_mlogloss:2.35956 eval_mlogloss:2.39674
170] train_mlogloss:2.35956 eval_mlogloss:2.39674
171] train_mlogloss:2.35956 eval_mlogloss:2.39678
172] return(train_mlogloss, eval_mlogloss)
173] train_mlogloss:2.35851 eval_mlogloss:2.39672
174] train_mlogloss:2.35834 eval_mlogloss:2.3967
175] train_mlogloss:2.35814 eval_mlogloss:2.39672
176] train_mlogloss:2.35814 eval_mlogloss:2.39669
177] labels=train_labels; eval_labels=eval_labels; train_mlogloss:2.35769 eval_mlogloss:2.39669
178] train_mlogloss:2.35769 eval_mlogloss:2.39669
179] train_mlogloss:2.35769 eval_mlogloss:2.39669
180] train_mlogloss:2.35738 eval_mlogloss:2.39665
181] plt.figure(figsize=(10,7)) eval_mlogloss:2.39668
182] for train_mlogloss in range(1,10): eval_mlogloss:2.39667
183] plt.scatter(train_labels, eval_labels, s=100, c='r', label='train')
184] train_mlogloss:2.35681 eval_mlogloss:2.39667
185] train_mlogloss:2.35681 eval_mlogloss:2.39667
186] train_mlogloss:2.35681 eval_mlogloss:2.39667
187] train_mlogloss:2.35681 eval_mlogloss:2.39672
188] train_mlogloss:2.35681 eval_mlogloss:2.39672
189] train_mlogloss:2.35681 eval_mlogloss:2.3967
190] plt.xlabel('log_loss') eval_mlogloss:2.39668
191] plt.ylabel('log_loss') eval_mlogloss:2.39667
192] plt.title('log_loss') eval_mlogloss:2.39665
193] plt.legend(log_loss) eval_mlogloss:2.39668
194] plt.savefig('log_loss.png') eval_mlogloss:2.39669
195] plt.show() eval_mlogloss:2.39669
196] train_mlogloss:2.3551 eval_mlogloss:2.39671
197] plt.figure(figsize=(10,7)) eval_mlogloss:2.39676
198] for train_mlogloss in range(1,10): eval_mlogloss:2.39675
199] plt.scatter(train_labels, eval_labels, s=100, c='r', label='train')
200] train_mlogloss:2.35456 eval_mlogloss:2.39679
201] train_mlogloss:2.35456 eval_mlogloss:2.39681
202] train_mlogloss:2.35456 eval_mlogloss:2.39683
203] train_mlogloss:2.35456 eval_mlogloss:2.39686
204] train_mlogloss:2.35456 eval_mlogloss:2.39683
205] train_mlogloss:2.35456 eval_mlogloss:2.39683
206] plt.xlabel('log_loss') eval_mlogloss:2.39683
207] plt.ylabel('log_loss') eval_mlogloss:2.39679
208] plt.title('log_loss') eval_mlogloss:2.39681
209] plt.legend(log_loss) eval_mlogloss:2.39682
210] plt.savefig('log_loss.png') eval_mlogloss:2.39685
211] plt.show() eval_mlogloss:2.39688
212] train_mlogloss:2.35293 eval_mlogloss:2.39695

```

```

[213] plt.figure(figsize=(10, 5)) eval-mlogloss: 2.39693
[214] for train_mlogloss in range(1, 10):
[215]     plt.scatter(train_mlogloss, np.where(train_mlogloss < 2.39693), 1),
[216]     train_mlogloss, np.where(train_mlogloss < 2.39693), 2],
[217]     train_mlogloss, eval-mlogloss: 2.39696
[218]     train_mlogloss, eval-mlogloss: 2.39696
[219]     train_mlogloss, eval-mlogloss: 2.39695
[220]     train_mlogloss, eval-mlogloss: 2.39692
[221]     train_mlogloss, eval-mlogloss: 2.3969
[222] plt.xlabel('logloss') eval-mlogloss: 2.39691
[223] plt.ylabel('logloss') eval-mlogloss: 2.3969
[224] plt.title('logloss') eval-mlogloss: 2.39688
[225] plt.legend(['logloss']) eval-mlogloss: 2.39692
[226] plt.savefig('logloss.png') eval-mlogloss: 2.39691
[227] plt.show() eval-mlogloss: 2.39689
[228] train_mlogloss: 2.35086 eval-mlogloss: 2.3969
[229] train_mlogloss: 2.35074 eval-mlogloss: 2.39691
[230] map_train_mlogloss: 2.3506 eval-mlogloss: 2.39689
[231] labels = sorted(unique(train_mlogloss)) eval-mlogloss: 2.39689
[232] mapping = dict(zip(labels, range(len(labels)))) eval-mlogloss: 2.39696
[233] for train_mlogloss in range(1, 10): eval-mlogloss: 2.39699
[234]     mapping[train_mlogloss] = i eval-mlogloss: 2.39699
[235] train_mlogloss.replace(mapping[train_mlogloss]) eval-mlogloss: 2.39703
[236] return train_mlogloss eval-mlogloss: 2.39706
[237] train_mlogloss: 2.34971 eval-mlogloss: 2.39705
[238] train_mlogloss: 2.34959 eval-mlogloss: 2.39707
[239] read_train_mlogloss: 2.34945 eval-mlogloss: 2.39705
[240] # App events: 2.34931 eval-mlogloss: 2.39705
[241] print('App events: 2.34931 eval-mlogloss: 2.39703')
[242] # App events: 2.34931 eval-mlogloss: 2.39702
Stopping: 2.34931 eval-mlogloss: 2.39702
[192] # Active users: 2.34931 eval-mlogloss: 2.39702
ape['active'] = ape.groupby(
Validating event_id')[ 'is_active'].transform('sum')
Predape['is_installed', 'is_active'], axis=1, inplace=True)
Trainape['is_installed', 'is_active'], axis=1, inplace=True)
LS: 2.34931 eval-mlogloss: 2.39702
Writing submission: submission_2.39667065647_2017-12-07-17-58.csv
# Events
print('Reading events...')
events = pd.read_csv('data/events.csv', dtype={'device_id': np.str})
events['counts'] = events.groupby(
['device_id'])['event_id'].transform('count')

print('Making events features...')
# The idea here is to count the number of installed apps using the data
# from app_events.csv above. Also to count the number of active apps.
events = pd.merge(events, ape, how='left', on='event_id', left_index=True)

# Below is the original events_small table
# events_small = events[['device_id', 'counts']].drop_duplicates('device_id', keep='first')
# And this is the new events_small table with two extra features
events_small = events[['device_id', 'counts', 'installed',

```

```

        'active']] .drop_duplicates('device_id',
                                   keep='first')

# Phone brand
print('Reading phone brands...')
pbd = pd.read_csv('data/phone_brand_device_model.csv',
                  dtype={'device_id': np.str})
pbd.drop_duplicates('device_id', keep='first', inplace=True)
pbd = map_column(pbd, 'phone_brand')
pbd = map_column(pbd, 'device_model')

# Train
print('Reading train data...')
train = pd.read_csv('data/gender_age_train.csv',
                    dtype={'device_id': np.str})
train = map_column(train, 'group')
train = train.drop(['age'], axis=1)
train = train.drop(['gender'], axis=1)
print('Merging features with train data...')
train = pd.merge(train, pbd, how='left', on='device_id', left_index=
True)
train = pd.merge(train,
                  events_small,
                  how='left',
                  on='device_id',
                  left_index=True)
train.fillna(-1, inplace=True)

# Test
print('Reading test data...')
test = pd.read_csv('data/gender_age_test.csv',
                   dtype={'device_id': np.str})
print('Merging features with test data...\n')
test = pd.merge(test, pbd, how='left', on='device_id', left_index=Tr
ue)
test = pd.merge(test,
                 events_small,
                 how='left',
                 on='device_id',
                 left_index=True)
test.fillna(-1, inplace=True)

# Features
features = list(test.columns.values)
features.remove('device_id')
return train, test, features

train, test, features = read_train_test()
print('Length of train: ', len(train))
print('Length of test: ', len(test))
print('Features [{ }]: { }\n'.format(len(features), sorted(features)))
train_df = pd.DataFrame(data=train)
X = train_df.drop(['group', 'device_id'], axis=1).values
Y = train_df['group'].values
tsne_data, tsne_groups = run_tsne(X, Y)
tsne_vis(tsne_data, tsne_groups)

```



```
Reading app events...
Reading events...
Making events features...
Reading phone brands...
Reading train data...
Merging features with train data...
Reading test data...
Merging features with test data...

Length of train: 74645
Length of test: 112071
Features [5]: ['active', 'counts', 'device_model', 'installed', 'phone_brand']

[t-SNE] Computing 91 nearest neighbors...
[t-SNE] Indexed 7465 samples in 0.024s...
[t-SNE] Computed neighbors for 7465 samples in 0.180s...
[t-SNE] Computed conditional probabilities for sample 1000 / 7465
[t-SNE] Computed conditional probabilities for sample 2000 / 7465
[t-SNE] Computed conditional probabilities for sample 3000 / 7465
[t-SNE] Computed conditional probabilities for sample 4000 / 7465
[t-SNE] Computed conditional probabilities for sample 5000 / 7465
[t-SNE] Computed conditional probabilities for sample 6000 / 7465
[t-SNE] Computed conditional probabilities for sample 7000 / 7465
[t-SNE] Computed conditional probabilities for sample 7465 / 7465
[t-SNE] Mean sigma: 0.000000
[t-SNE] Computed conditional probabilities in 0.793s
[t-SNE] Iteration 50: error = 67.1231308, gradient norm = 0.0067935 (50 iterations in 17.044s)
[t-SNE] Iteration 100: error = 59.3737221, gradient norm = 0.0023253 (50 iterations in 13.042s)
[t-SNE] Iteration 150: error = 56.5824242, gradient norm = 0.0015247 (50 iterations in 12.257s)
[t-SNE] Iteration 200: error = 55.1853828, gradient norm = 0.0010795 (50 iterations in 12.140s)
[t-SNE] Iteration 250: error = 54.3001022, gradient norm = 0.0007644 (50 iterations in 12.089s)
[t-SNE] KL divergence after 250 iterations with early exaggeration: 54.300102
[t-SNE] Iteration 300: error = 1.1290563, gradient norm = 0.0007107 (50 iterations in 13.919s)
[t-SNE] Iteration 350: error = 0.6922646, gradient norm = 0.0003353 (50 iterations in 14.901s)
[t-SNE] Iteration 400: error = 0.4979769, gradient norm = 0.0001582 (50 iterations in 15.257s)
[t-SNE] Iteration 450: error = 0.3815096, gradient norm = 0.0000979 (50 iterations in 15.627s)
[t-SNE] Iteration 500: error = 0.3082786, gradient norm = 0.0000712 (50 iterations in 15.849s)
[t-SNE] Error after 500 iterations: 0.308279
```

